



# Training deep neural density estimators to identify mechanistic models of neural dynamics

Pedro J Gonçalves<sup>1,2†\*</sup>, Jan-Matthis Lueckmann<sup>1,2†\*</sup>, Michael Deistler<sup>1,3†\*</sup>,  
Marcel Nonnenmacher<sup>1,2,4</sup>, Kaan Öcal<sup>2,5</sup>, Giacomo Bassetto<sup>1,2</sup>,  
Chaitanya Chintaluri<sup>6,7</sup>, William F Podlaski<sup>6</sup>, Sara A Haddad<sup>8</sup>, Tim P Vogels<sup>6,7</sup>,  
David S Greenberg<sup>1,4</sup>, Jakob H Macke<sup>1,2,3,9\*</sup>

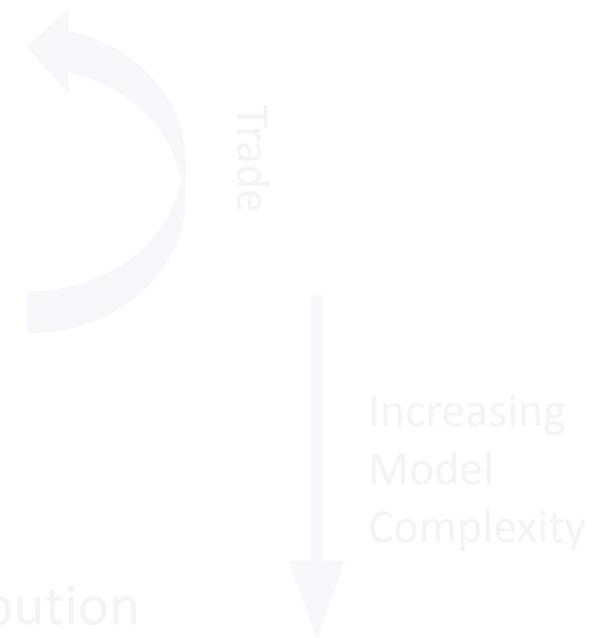
Computational and Systems Neuroscience Journal Club

11/3/2022

Abuzar Mahmood - Katz Lab

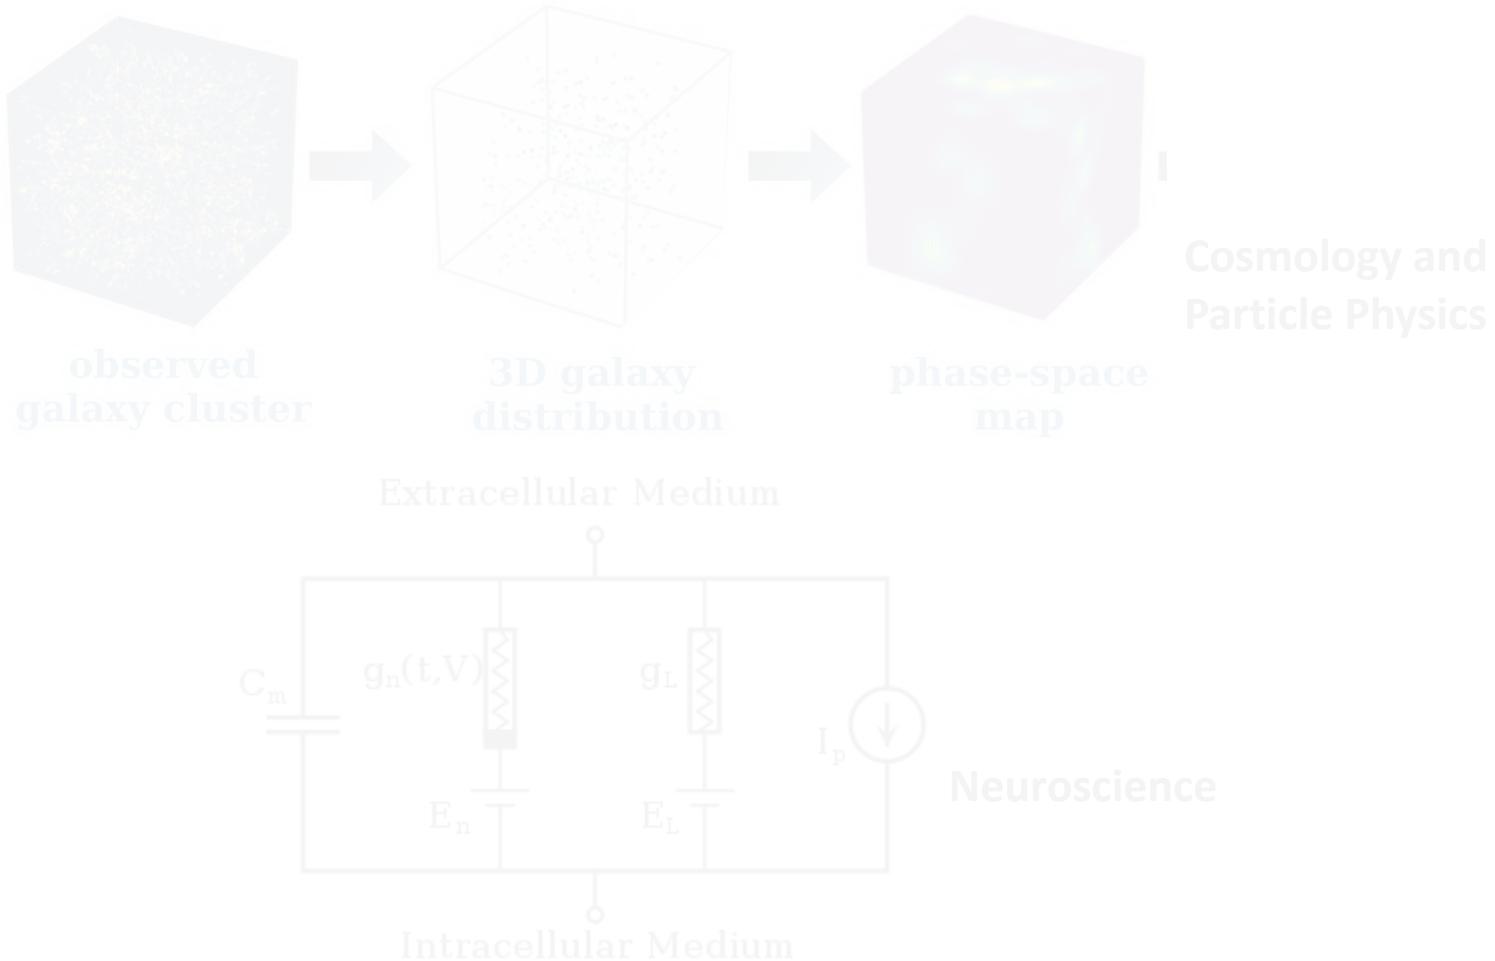
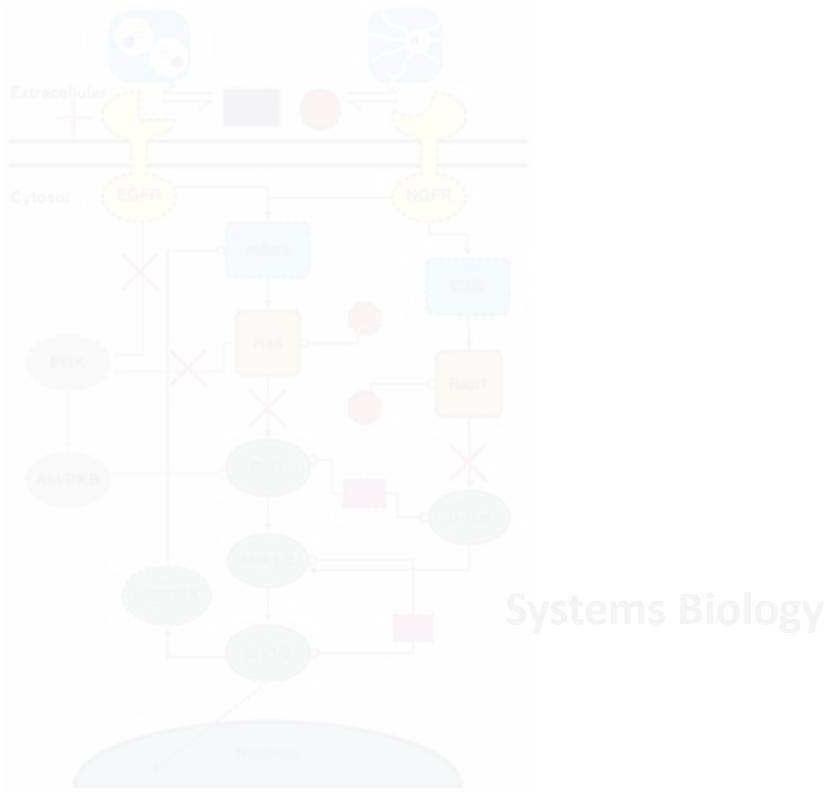
# Layout

- Motivation
- Caveats (what this talk will tell you, and what it won't)
- Background
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - Prior/other methods (epsilon / likelihood-free methods)
  - ABC and SMB-ABC
- Results
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. Inferring ion channel parameters
  5. Inferring parameters for Hodgkin-Huxley Model
  6. Inferring parameters for STG and analyzing posterior distribution



# Motivation - We need mechanistic models

- Does it work how we think it works?
- Model vs parameters?



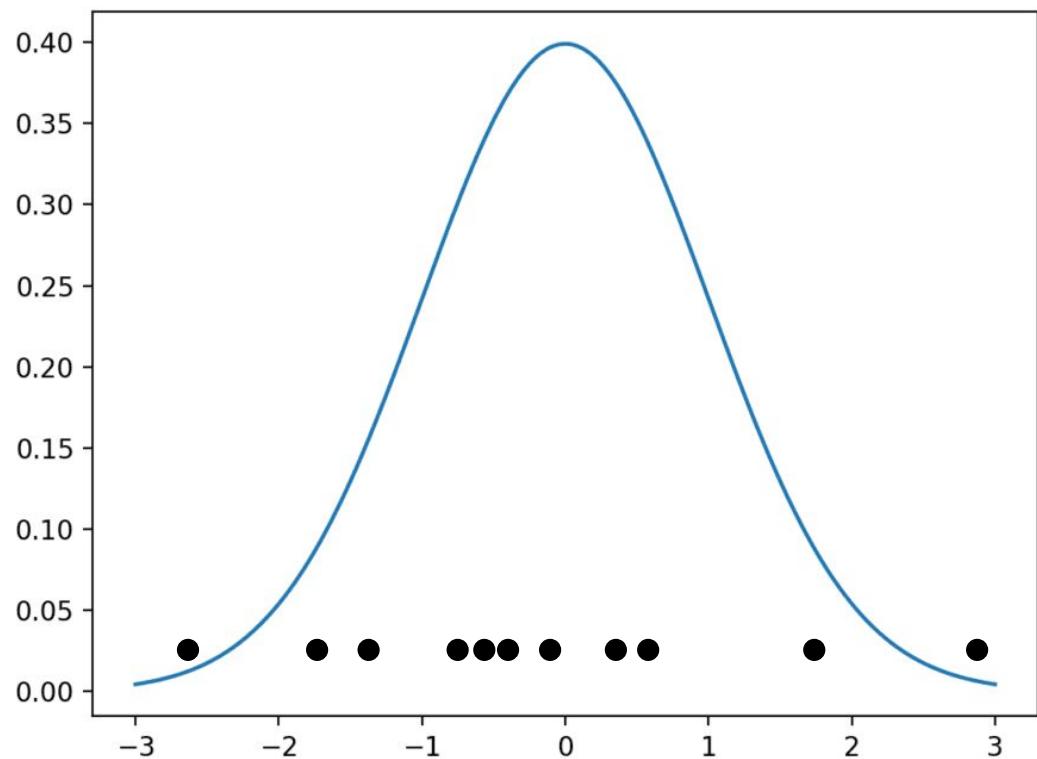
<https://dark.nbi.ku.dk/news/dark-news-21/simulation-based-inference-of-dynamical-galaxy-cluster-masses-with-3d-convolutional-neural-networks/>  
<https://sethna.lassp.cornell.edu/Sloppy/SloppySignalTransduction.html>

# Motivation - We need the distribution of relevant parameters

- Adjusting free parameters to be consistent with experimental observations
  - Essential for investigating whether the model agrees with reality
  - **For gaining insight into processes which cannot be measured experimentally (latent variables)**
- To understand how a model quantitatively explains data, necessary to find not only the best, but all parameter settings consistent with experimental observations
  - Something not possible/very difficult with point optimization techniques

# Motivation – Likelihood-free inference

- What is a likelihood?
  - How likely are data to come from a parametric distribution?



$$p(x_n) = \mathcal{N}(x|\mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

$$p(\{x_1, \dots, x_N\} | \pi_k, \mu_k, \sigma_k) = \prod_{n=1}^N p(x_n),$$

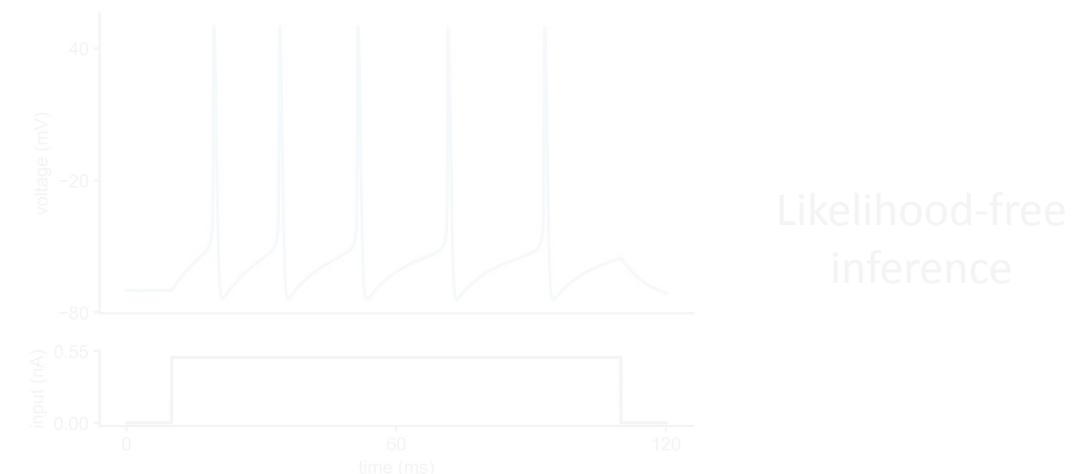
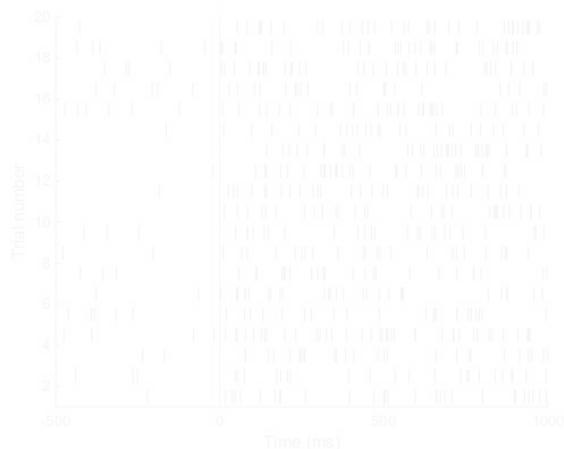
Likelihood allows us to use gradient-based methods to fit parameters

No likelihood  Not discussing likelihood-based methods

# Motivation – Current State of Affairs

- Gold standard for automated parameter identification is statistical inference, which uses the likelihood to quantify the match between parameters and data
- Likelihoods can be efficiently computed for purely statistical models commonly used in neuroscience, but are computationally intractable for most mechanistic models

Likelihood-based  
Inference  
(on parametric distributions)



# Layout

- Motivation
- **Caveats (what this talk will tell you, and what it won't)**
- Background
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - Related work (epsilon / likelihood-free methods)
    - ABC and SMB-ABC
- Results
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. ~~Inferring ion channel parameters~~
  5. Inferring parameters for Hodgkin-Huxley Model
  6. Inferring parameters for STG and analyzing posterior distribution

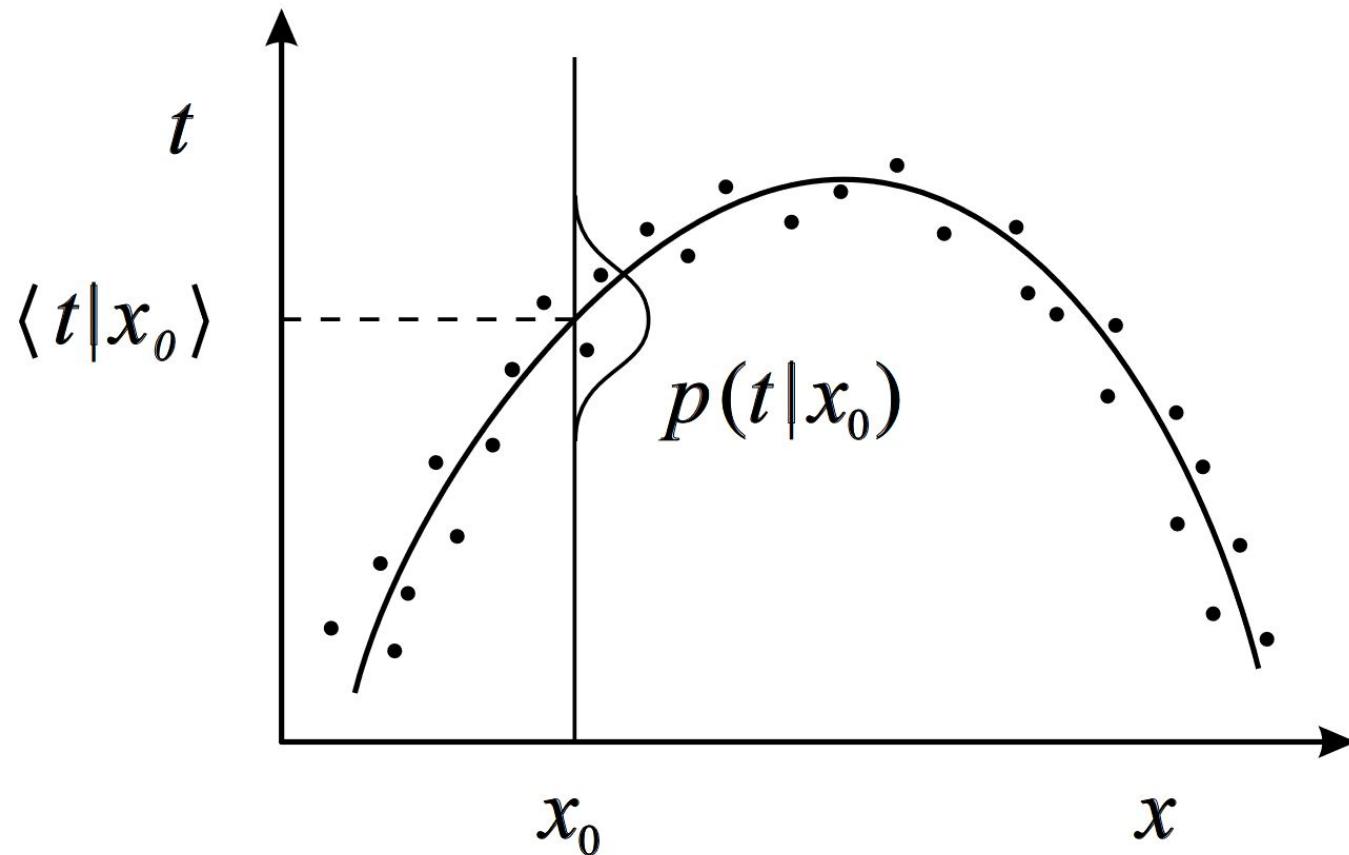
# Caveats

- Will not be diving into the math
- High-level concepts and results
- No new science :/

# Layout

- Motivation
- Caveats (what this talk will tell you, and what it won't)
- **Background**
  - **Conditional Density Estimation**
  - **Mixing Neural Networks with Probabilistic Models**
  - Related work (epsilon / likelihood-free methods)
    - ABC and SMB-ABC
- Results
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. ~~Inferring ion channel parameters~~
  5. Inferring parameters for Hodgkin-Huxley Model
  6. Inferring parameters for STG and analyzing posterior distribution

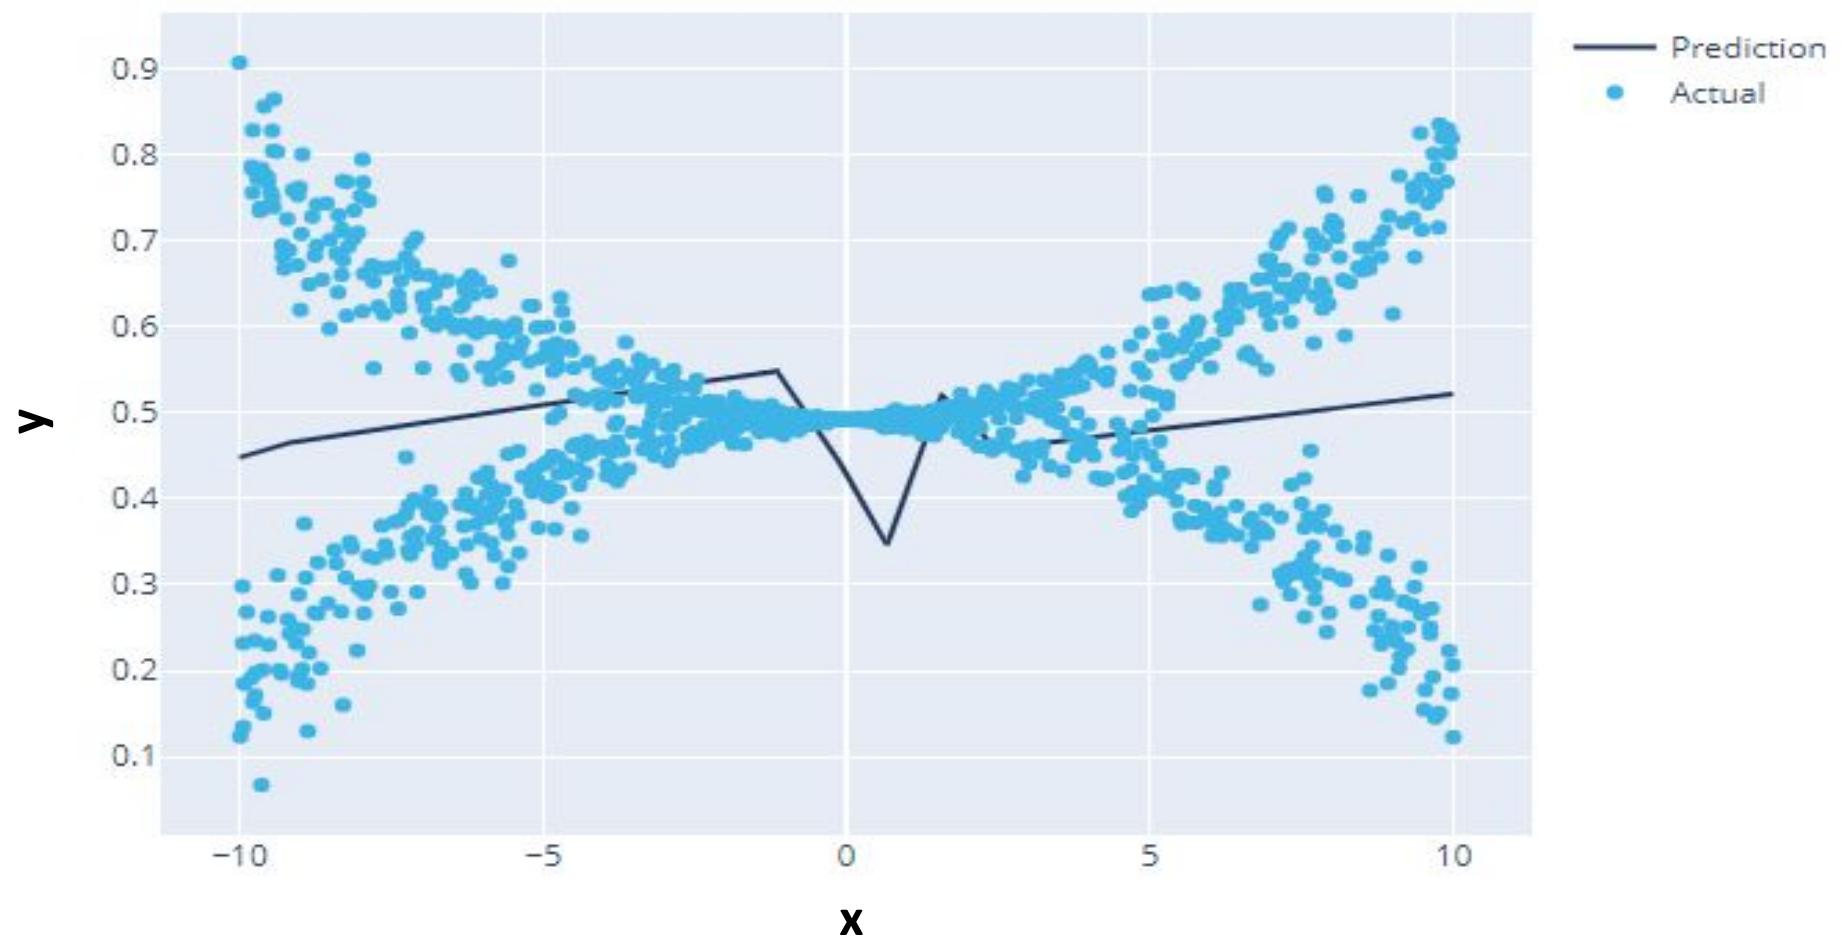
# Background – Conditional Density Estimation



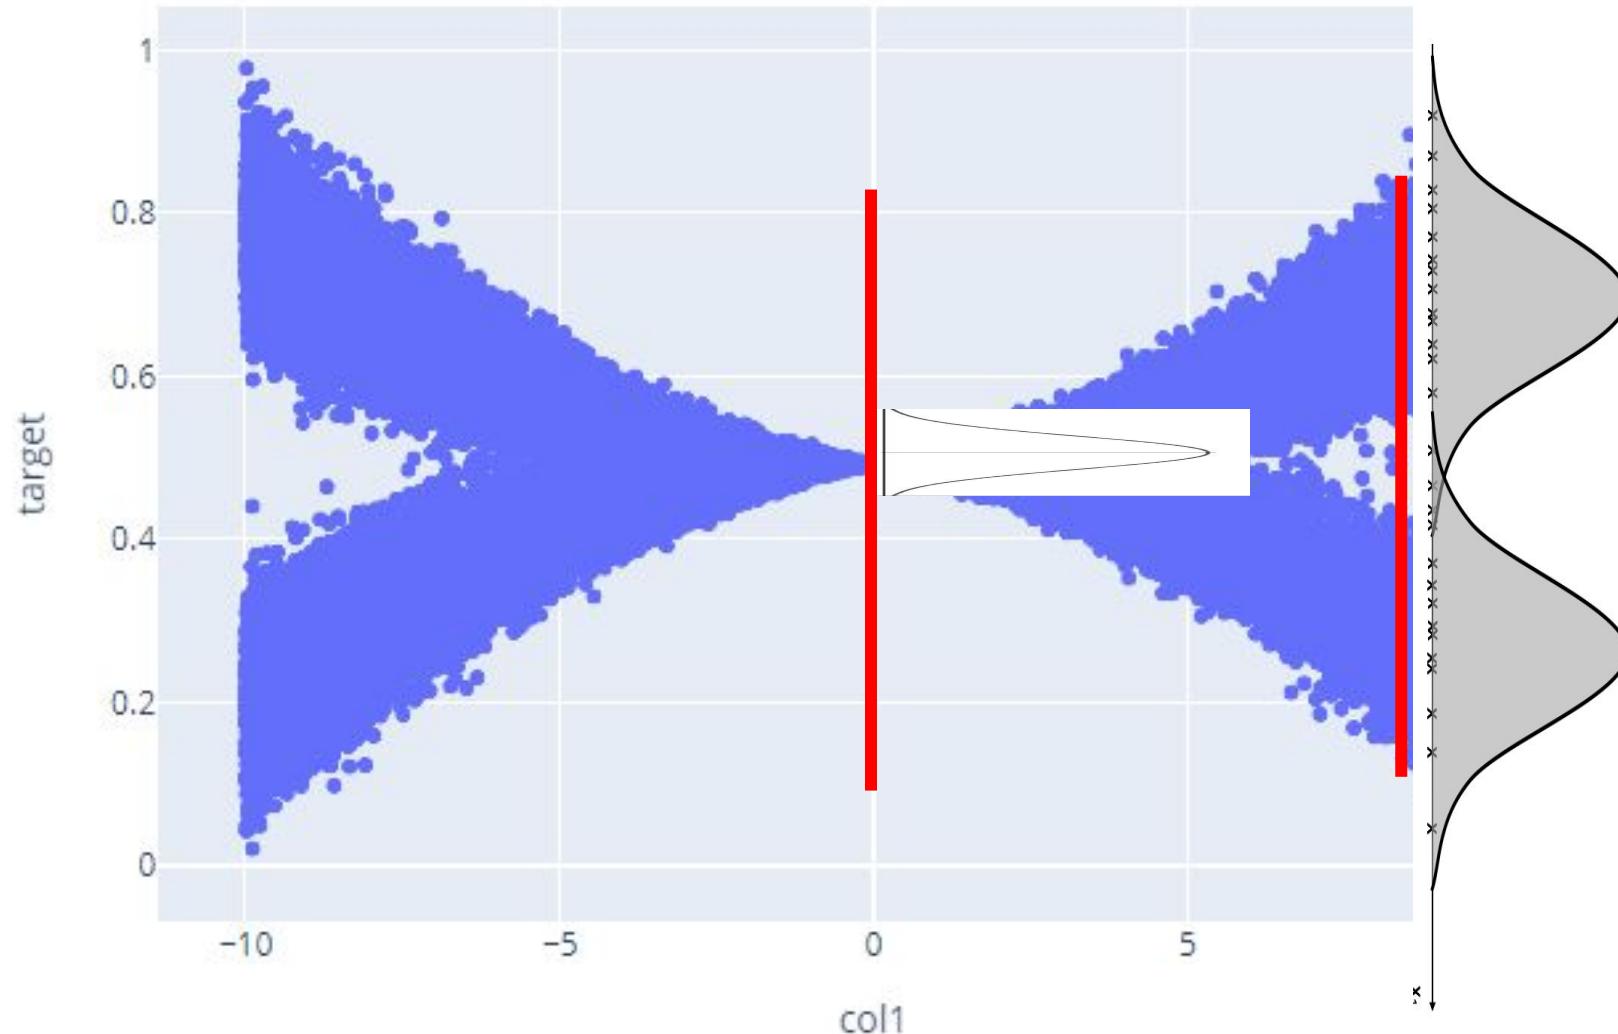
1. Why do we want to learn a conditional density?
  - Posterior  $\rightarrow p(\theta|x)$
2. Why do we need to learn a density rather than the mean value here?

# Background – Conditional Density Estimation

Vanilla Neural Network

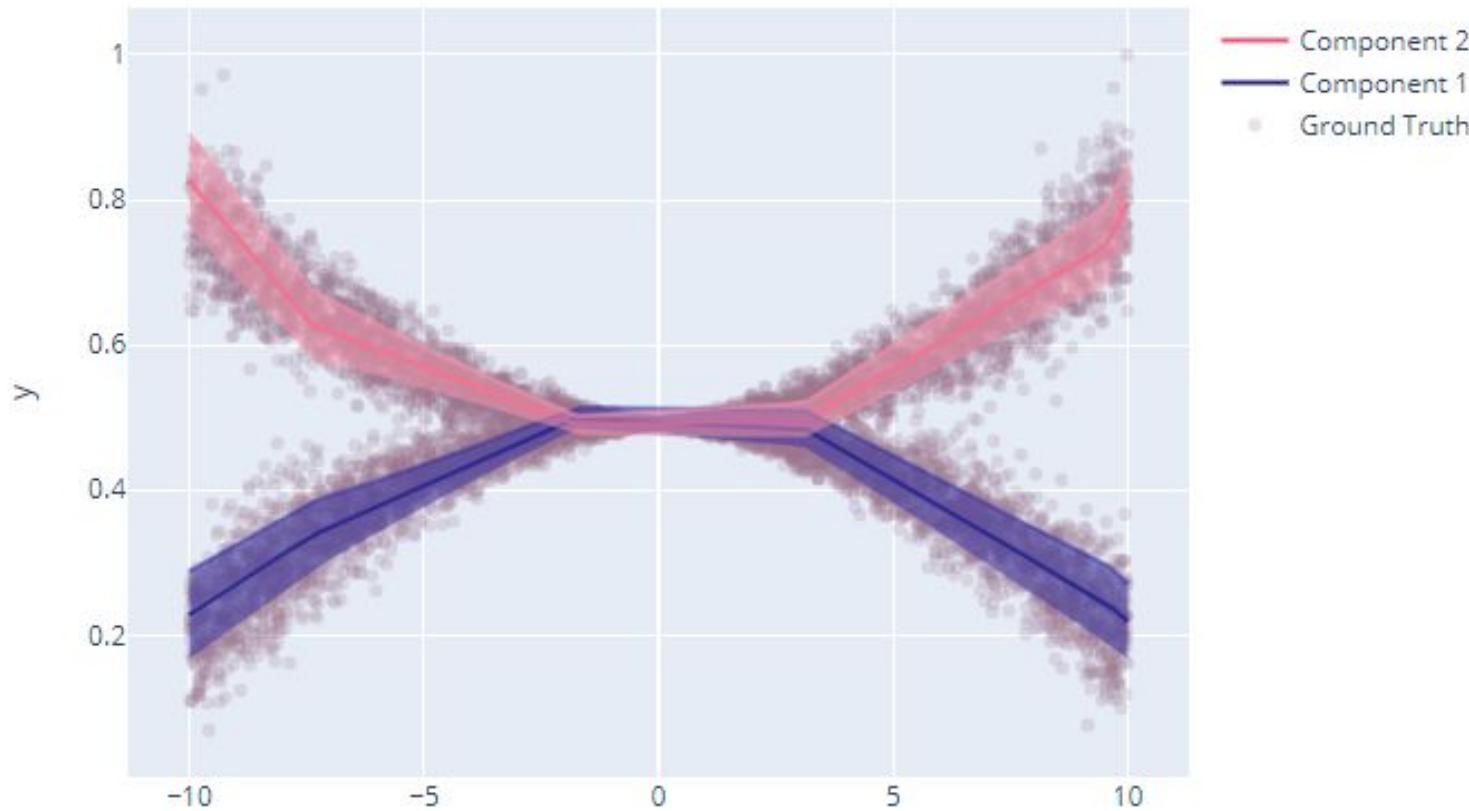


# Background – Conditional Density Estimation



# Background – Conditional Density Estimation

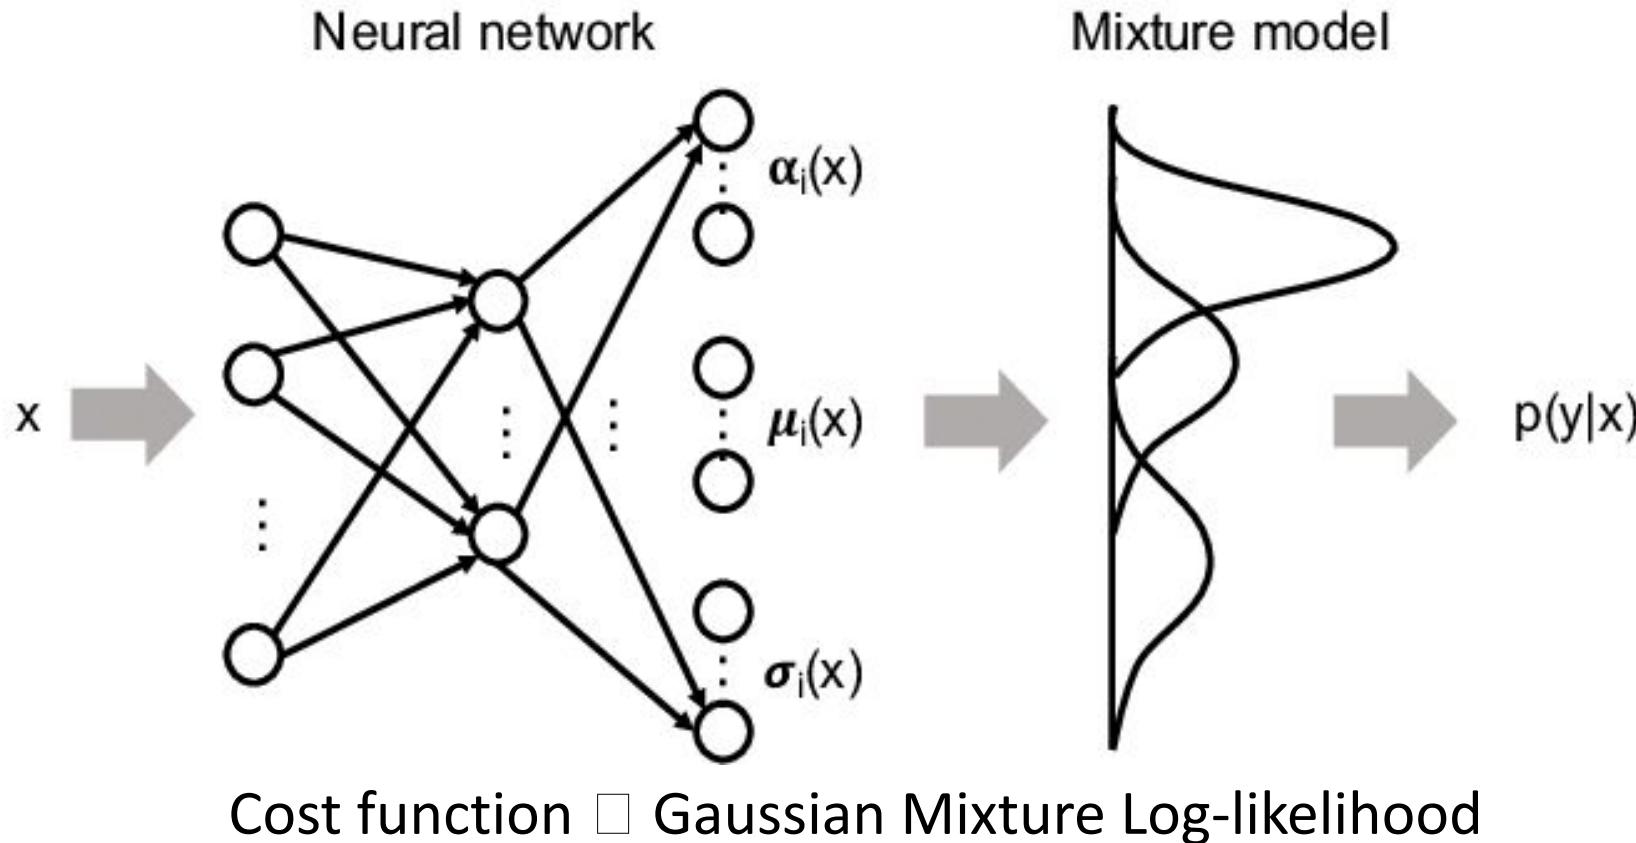
Mixture Density Network Prediction



1. How does one learn the conditional density?

# Background – Mixing Neural Networks with Probabilistic Models

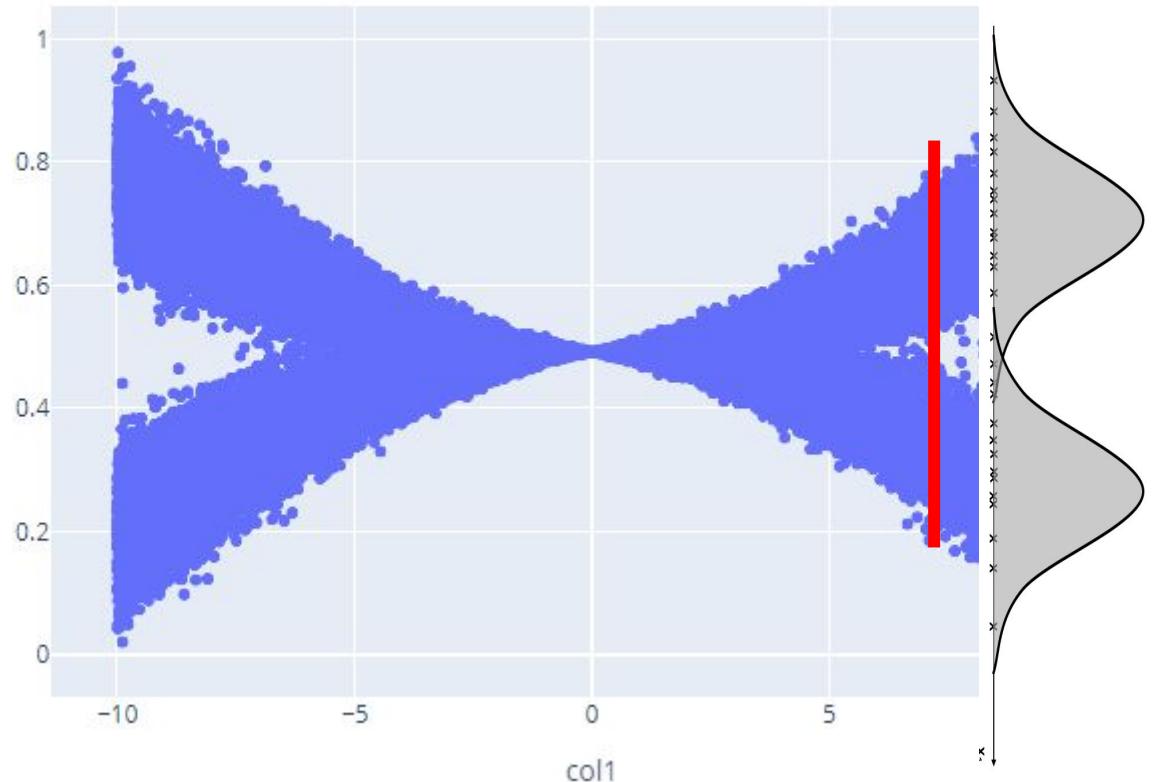
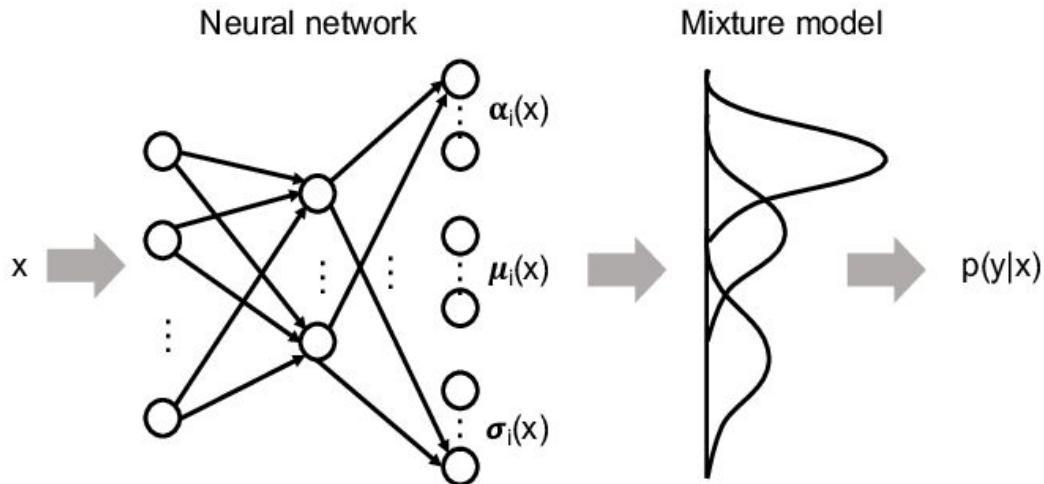
- Example – Mixture Density Network



- Are gaussian mixtures enough to learn complex distributions?
  - YES, given enough components
  - But authors also use normalizing flows

# Background – Mixing Neural Networks with Probabilistic Models

- Example – Mixture Density Network
  - **Differentiable representation of the density – Important for later**



# Layout

- Motivation
- Caveats (what this talk will tell you, and what it won't)
- **Background**
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - **Related work (epsilon / likelihood-free methods)**
    - **ABC and SMB-ABC**
- Results
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. ~~Inferring ion channel parameters~~
  5. Inferring parameters for Hodgkin-Huxley Model
  6. Inferring parameters for STG and analyzing posterior distribution

# Related work – Likelihood-free inference

- Approximate Bayesian Inference (ABC)

- Rejection sampling ABC ( $\epsilon$ )

1. Given model
2. Initialize random hyperparameters/parental variables from prior distributions
3. Generate sample
4. If  $| \text{target} - \text{sample} | < \epsilon$  (according to some distance metric), keep sample, else reject

- Problems

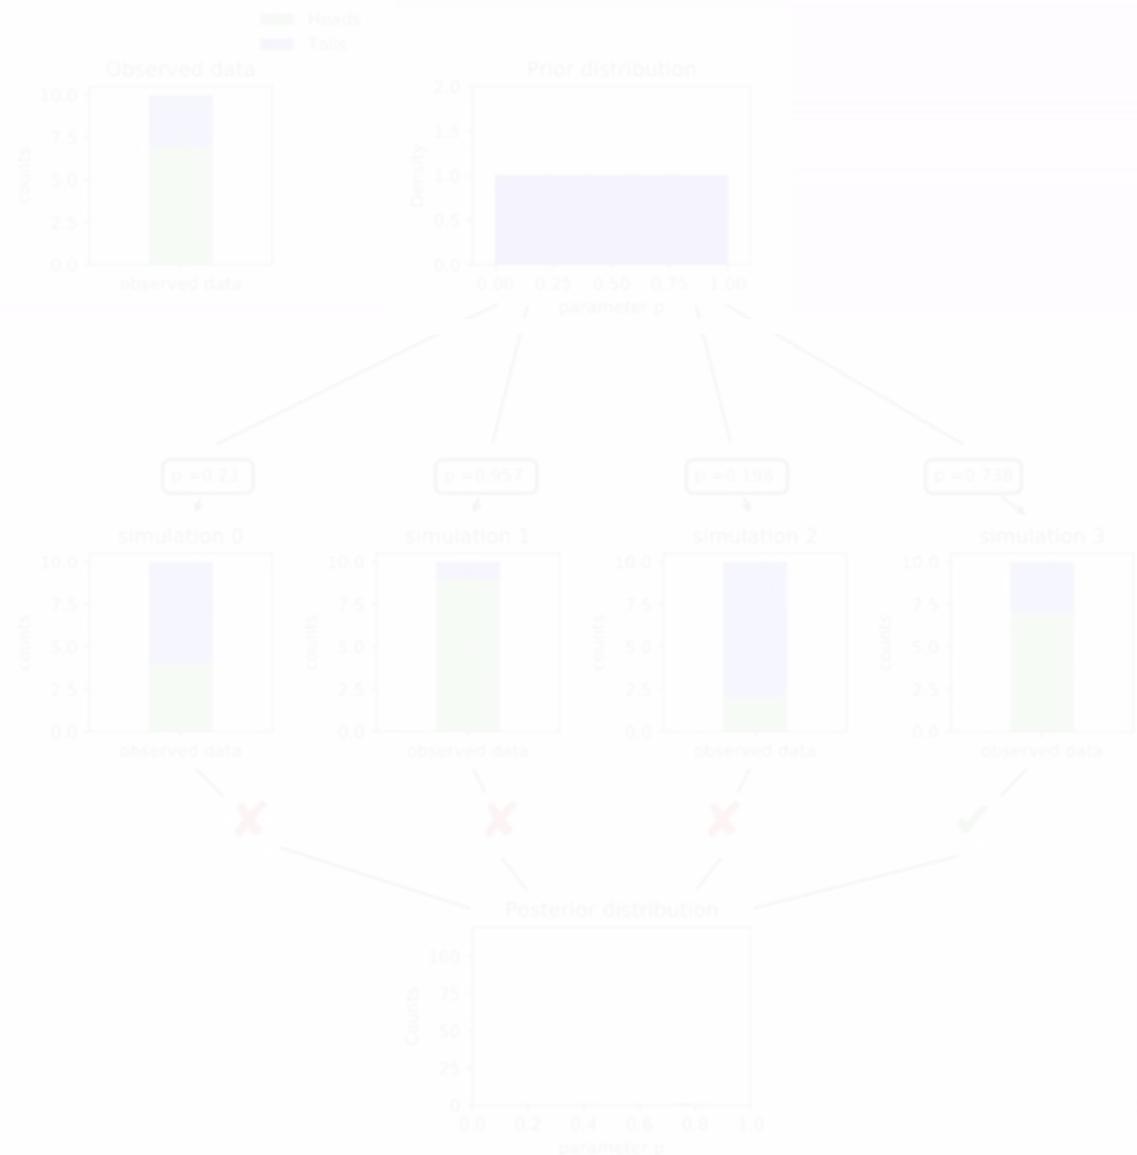
- Rejection rate too high for any interesting problem
- Sampling in areas of low likelihood

- Introduction to Approximate Bayesian Computation (ABC)

- <https://darrenjw.wordpress.com/2013/03/31/introduction-to-approximate-bayesian-computation-abc/>

- Tiny Data, Approximate Bayesian Computation and the Socks of Karl Broman

- <https://www.sumsar.net/blog/2014/10/tiny-data-and-the-socks-of-karl-broman/>



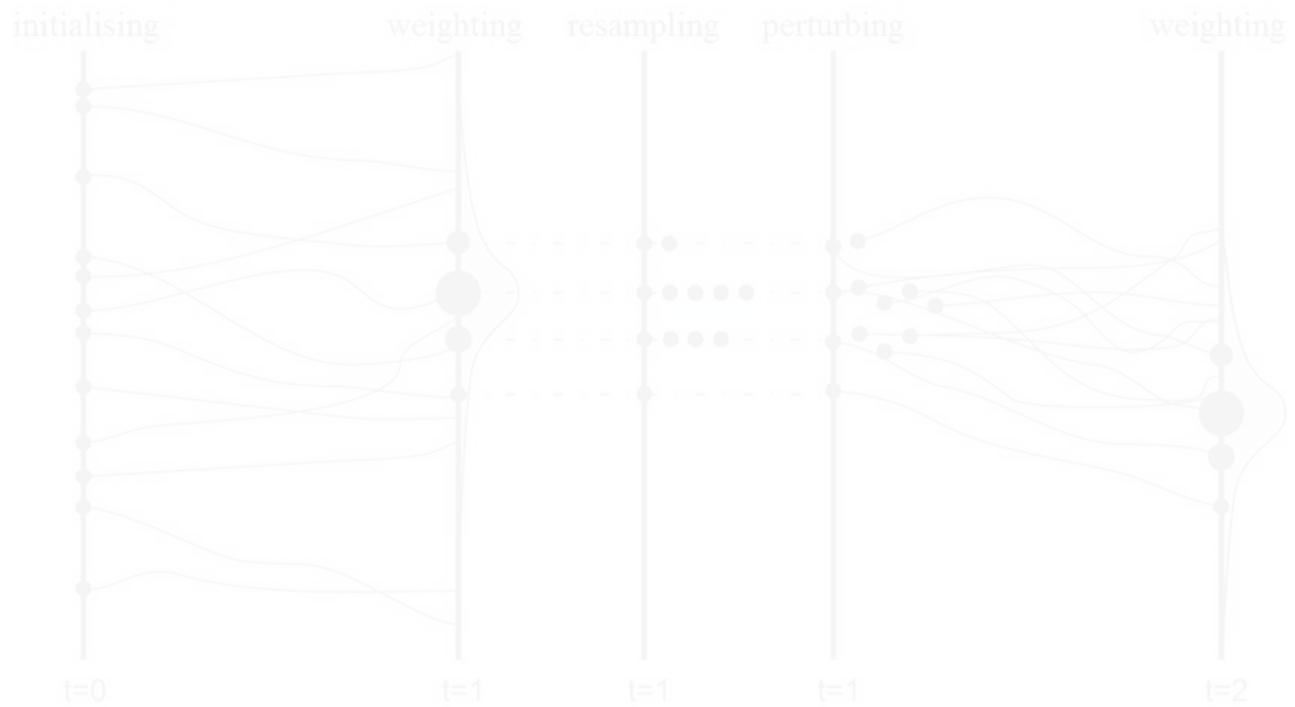
# Related work – Likelihood-free inference

- Sequential Monte-Carlo ABC

- Similar to regular SMC/Particle Filter, but with likelihood replaced with a cost/distance function

1. Given model
2. Initialize random hyperparameters/parental variables from prior distributions
3. Generate sample
4. If  $| \text{target} - \text{sample} | < \epsilon$  (according to some distance metric), keep sample, else reject
5. Kept samples + noise (from pre-specified kernel) become new prior distributions

- Like a genetic algorithm, but with distributions





# Training deep neural density estimators to identify mechanistic models of neural dynamics

Pedro J Gonçalves<sup>1,2†\*</sup>, Jan-Matthis Lueckmann<sup>1,2†\*</sup>, Michael Deistler<sup>1,3†\*</sup>,  
Marcel Nonnenmacher<sup>1,2,4</sup>, Kaan Öcal<sup>2,5</sup>, Giacomo Bassetto<sup>1,2</sup>,  
Chaitanya Chintaluri<sup>6,7</sup>, William F Podlaski<sup>6</sup>, Sara A Haddad<sup>8</sup>, Tim P Vogels<sup>6,7</sup>,  
David S Greenberg<sup>1,4</sup>, Jakob H Macke<sup>1,2,3,9\*</sup>

Part Deux

Computational and Systems Neuroscience Journal Club

11/10/2022

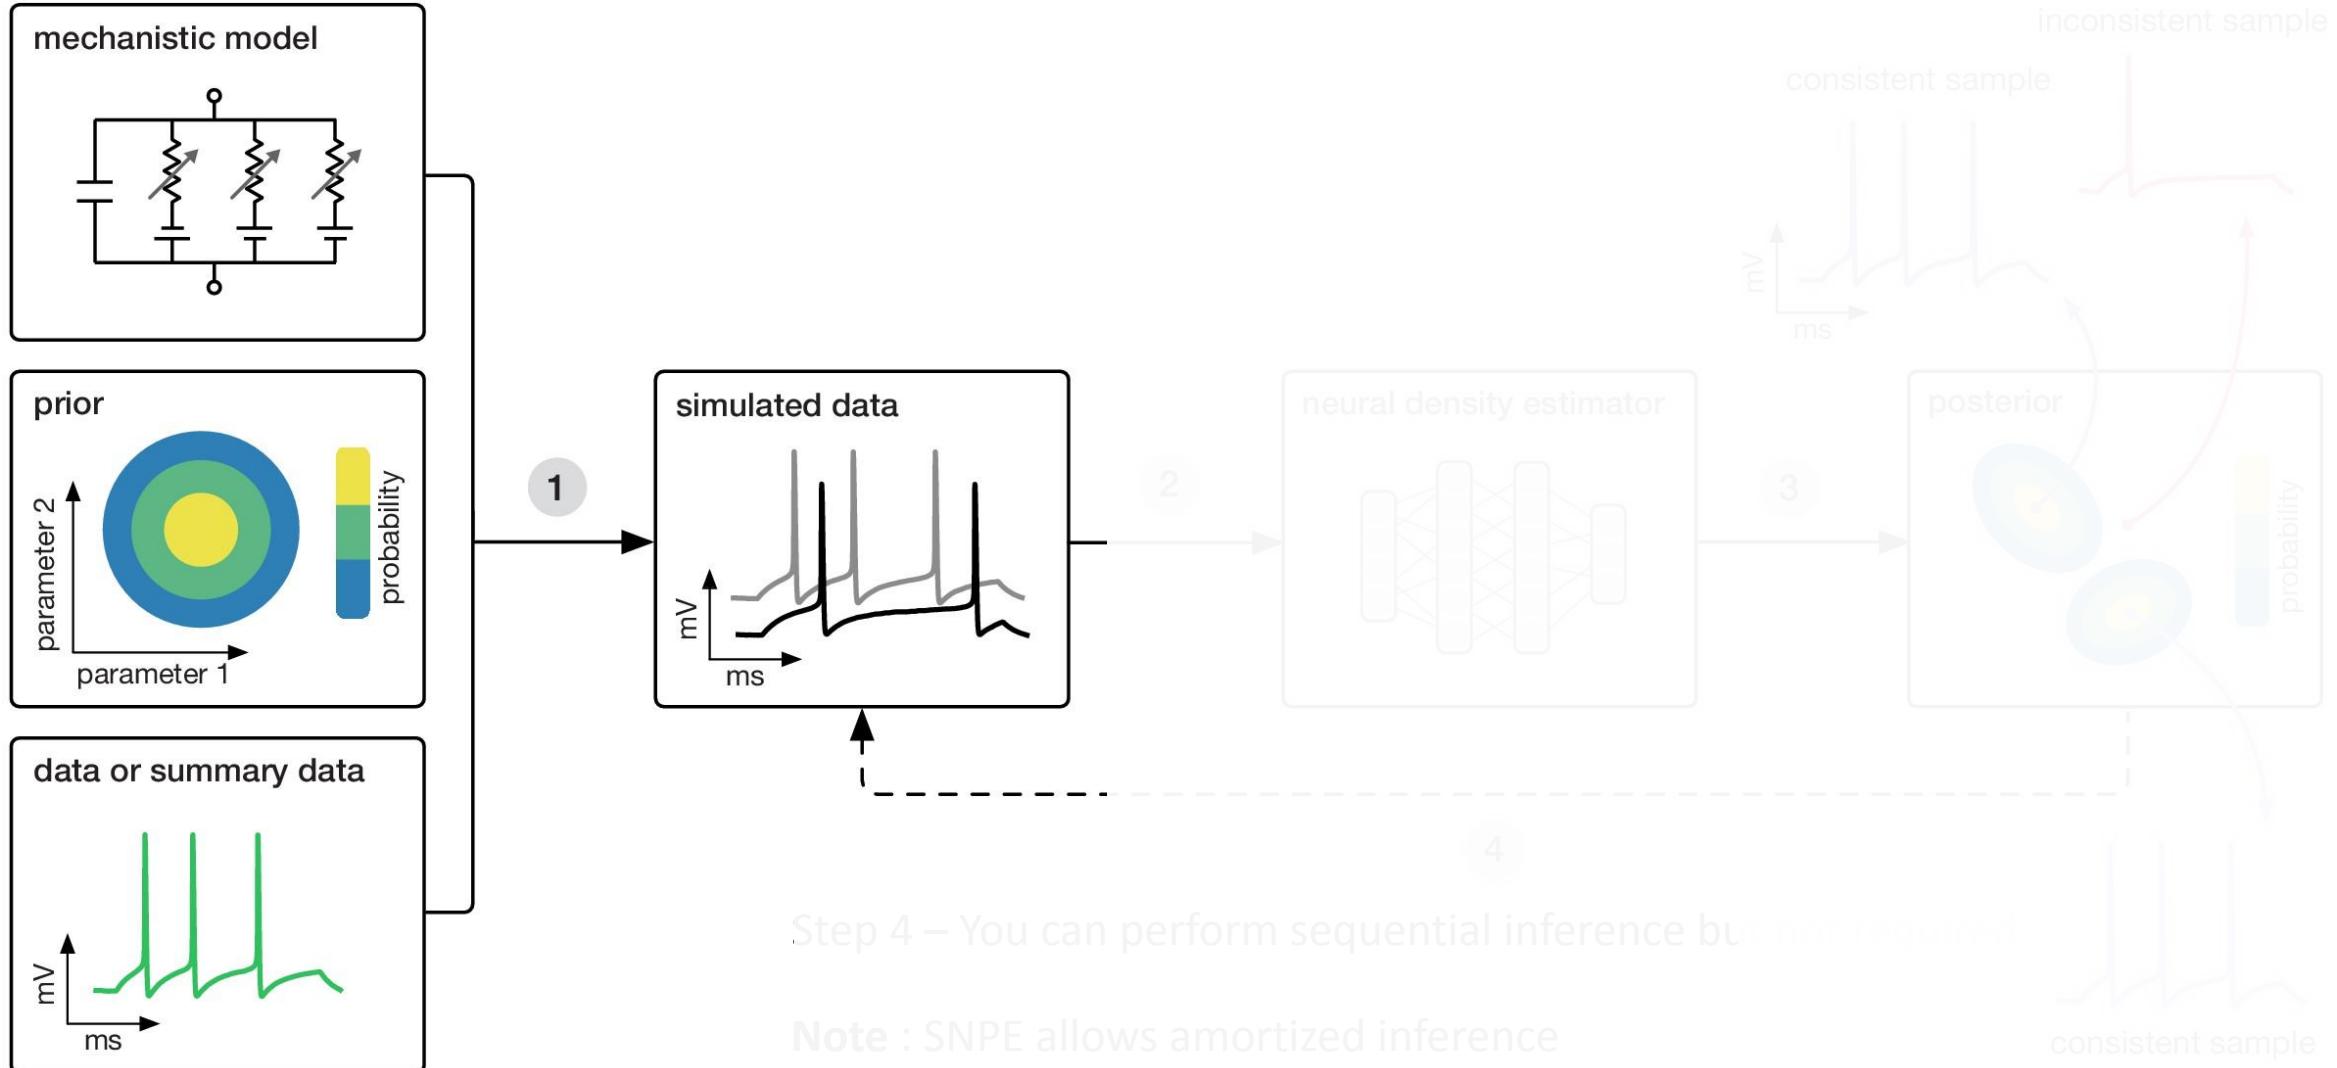
Abuzar Mahmood - Katz Lab

# Layout

- Motivation
- Caveats (what this talk will tell you, and what it won't)
- Background
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - Related work (epsilon / likelihood-free methods)
    - ABC and SMB-ABC
- **Results**
  1. **Recap of Methodology**
  2. **Comparison with established methodologies**
  3. **Inferring visual receptive field**
  4. ~~Inferring ion channel parameters~~
  5. Inferring parameters for Hodgkin-Huxley Model
  6. Inferring parameters for STG and analyzing posterior distribution

# Results – Recap of Methodology

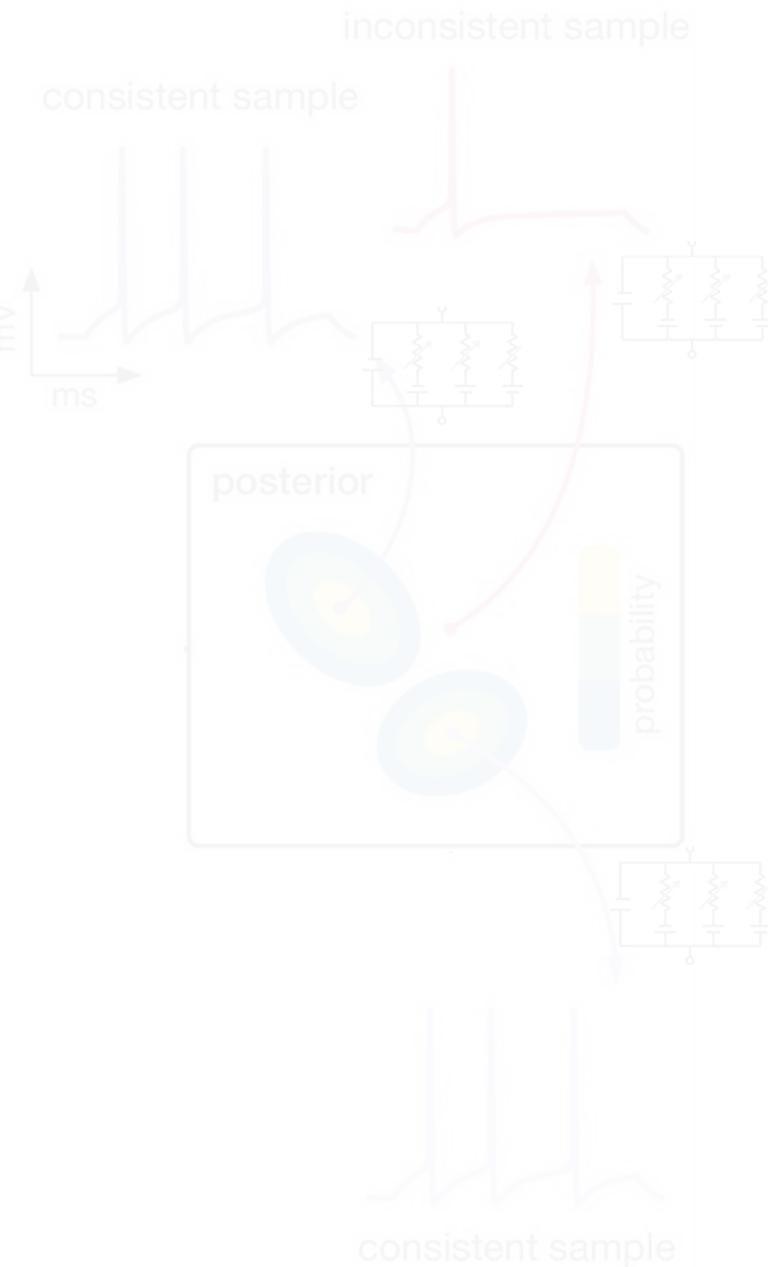
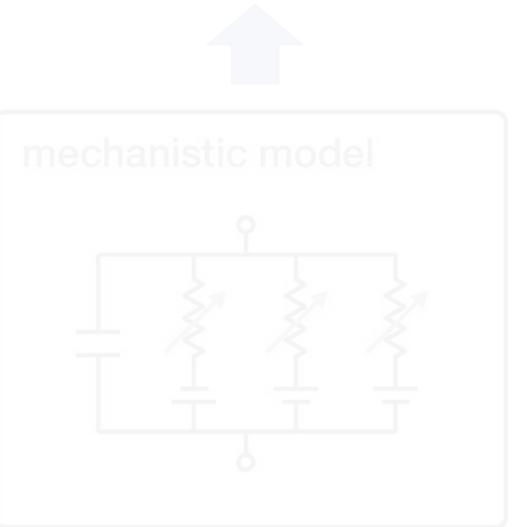
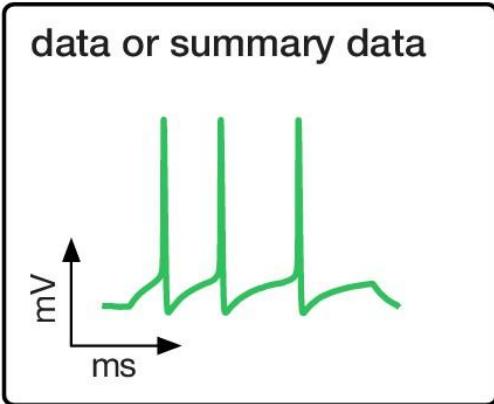
## SNPE – Sequential Neural Posterior Estimation



# Results – Recap of Methodology

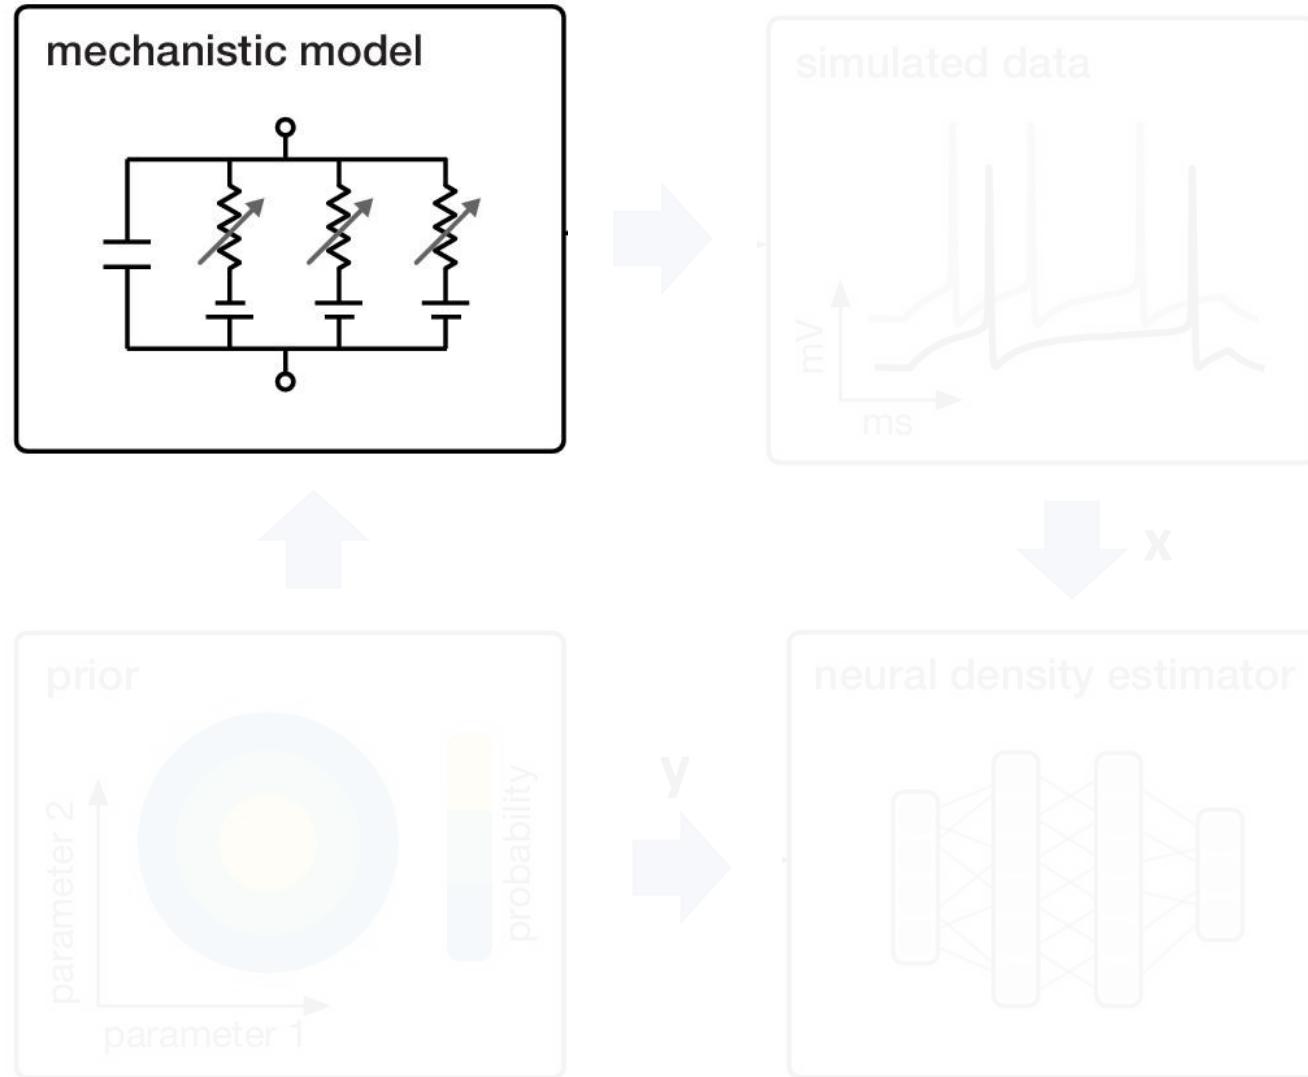
## The Goal

- Not prescribing utility of having the posterior distribution
- Distribution of parameters can provide similar outputs
- Having full distribution provides understanding of variance (uncertainty) and covariances



# Results – Recap of Methodology

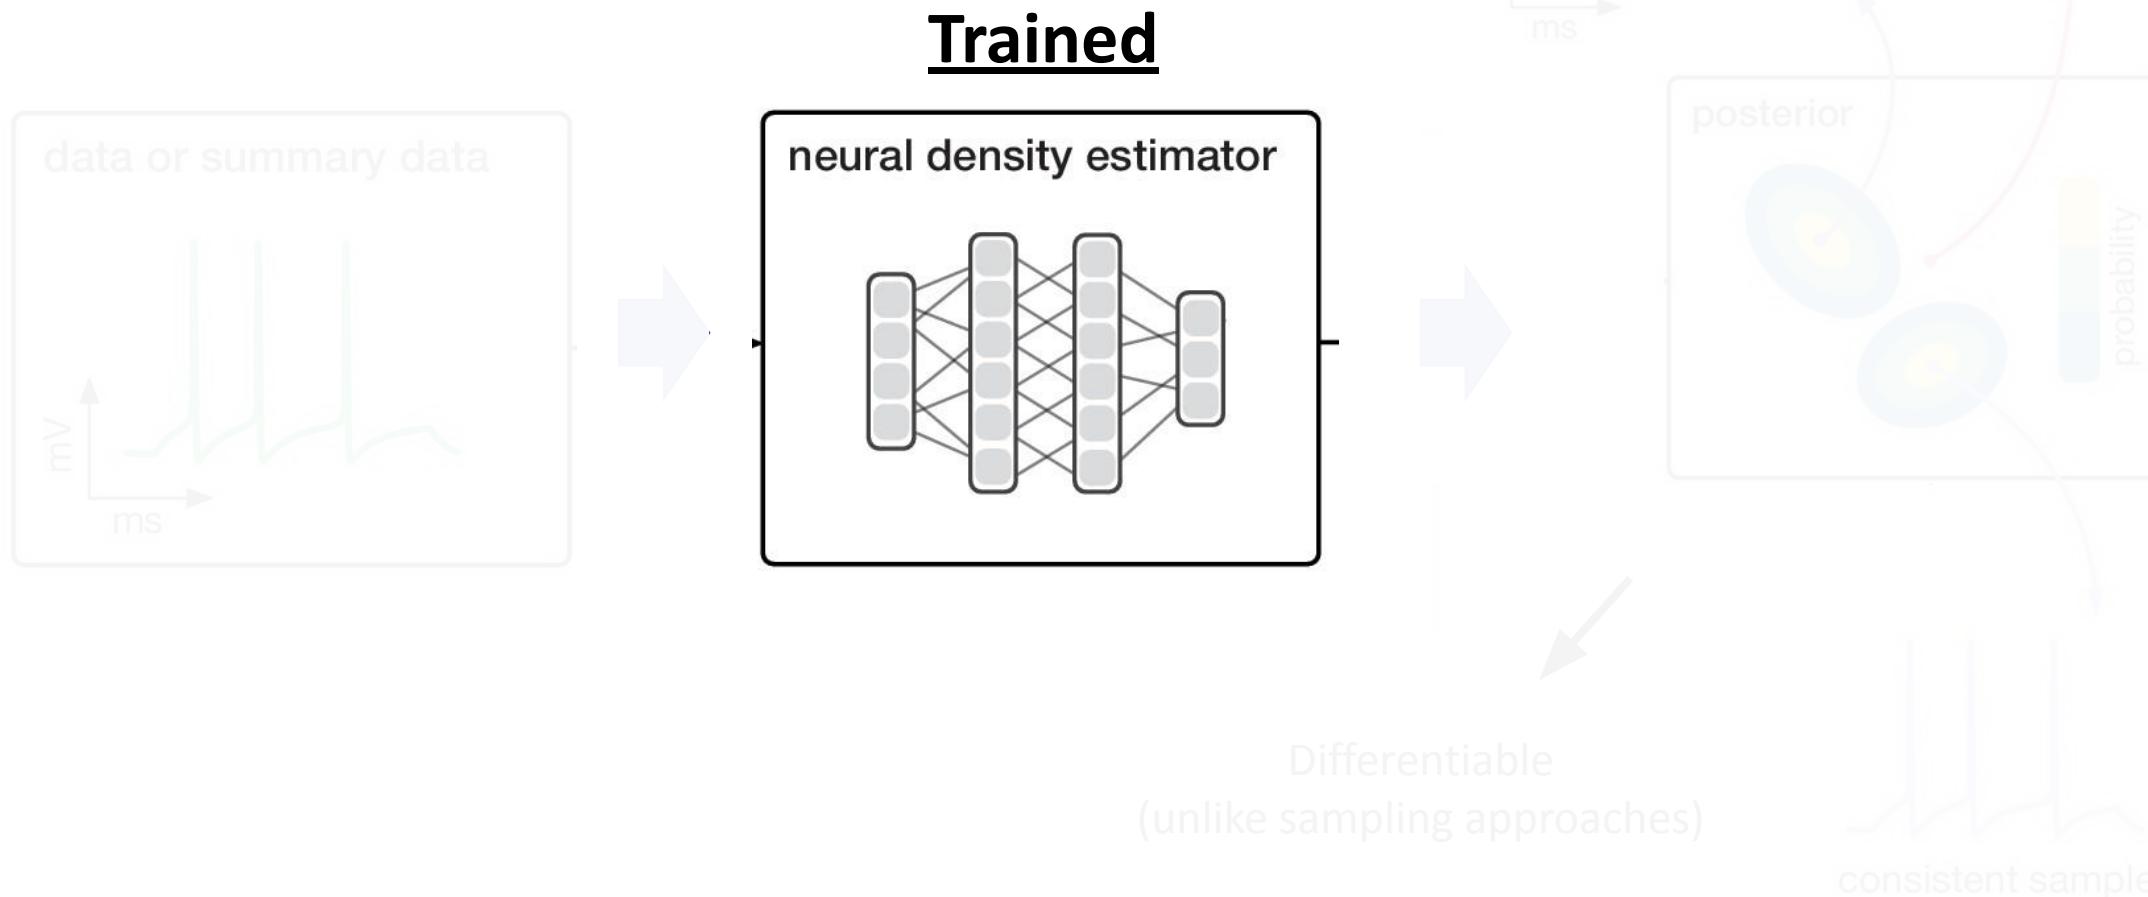
## SNPE – Sequential Neural Posterior Estimation



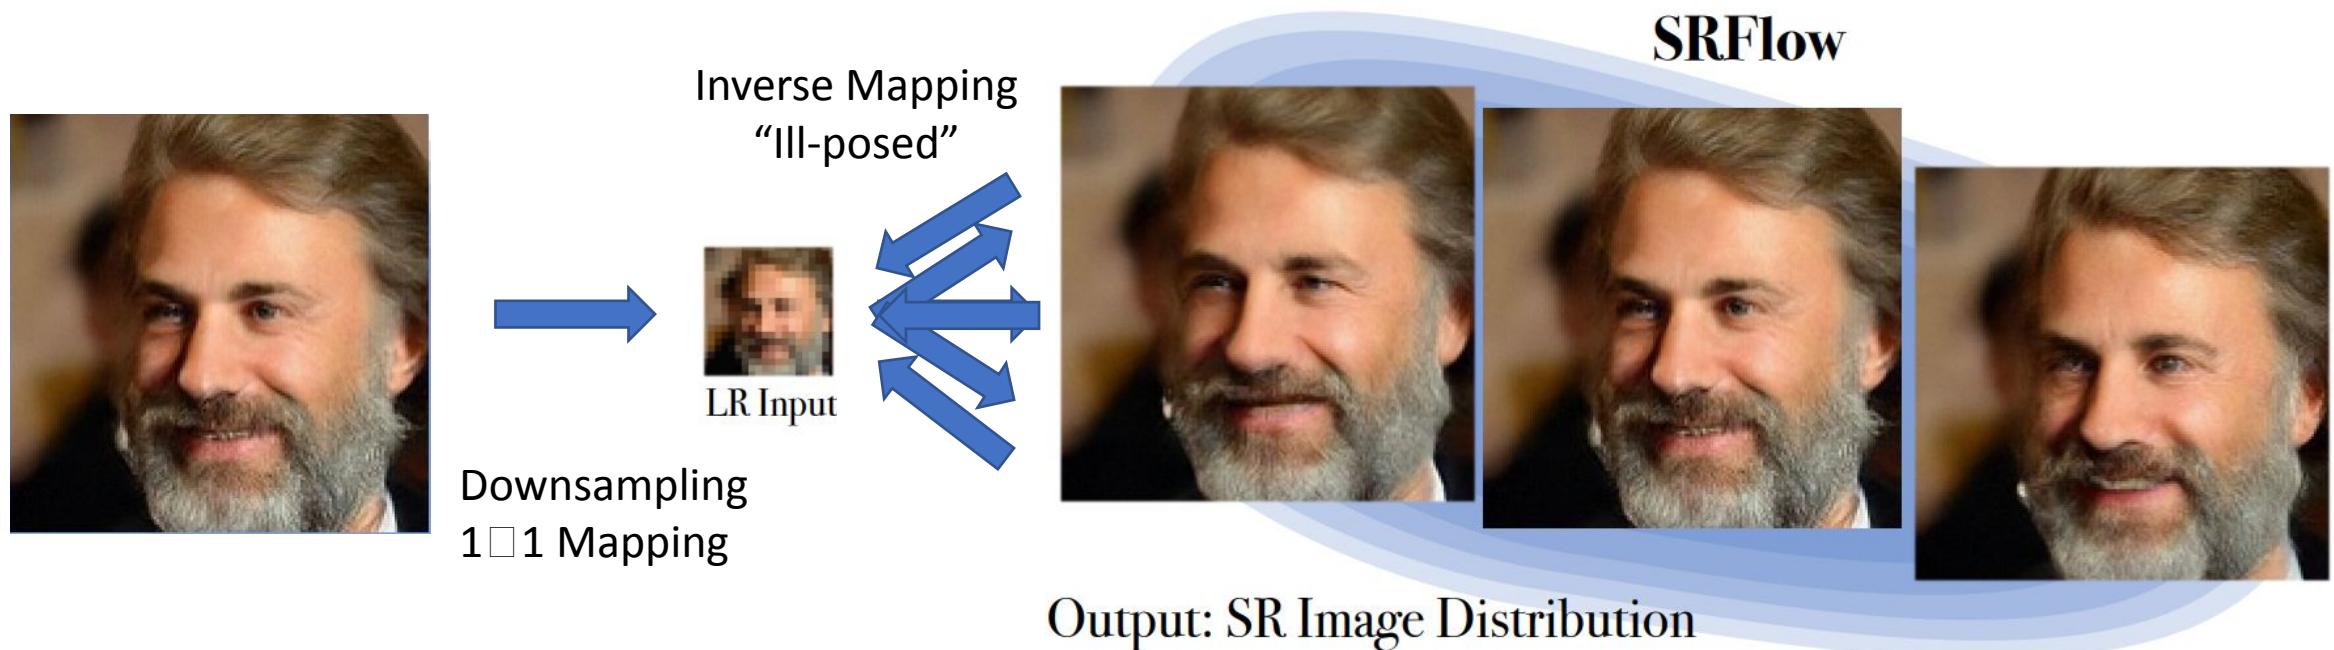
# Results – Recap of Methodology

## SNPE – Sequential Neural

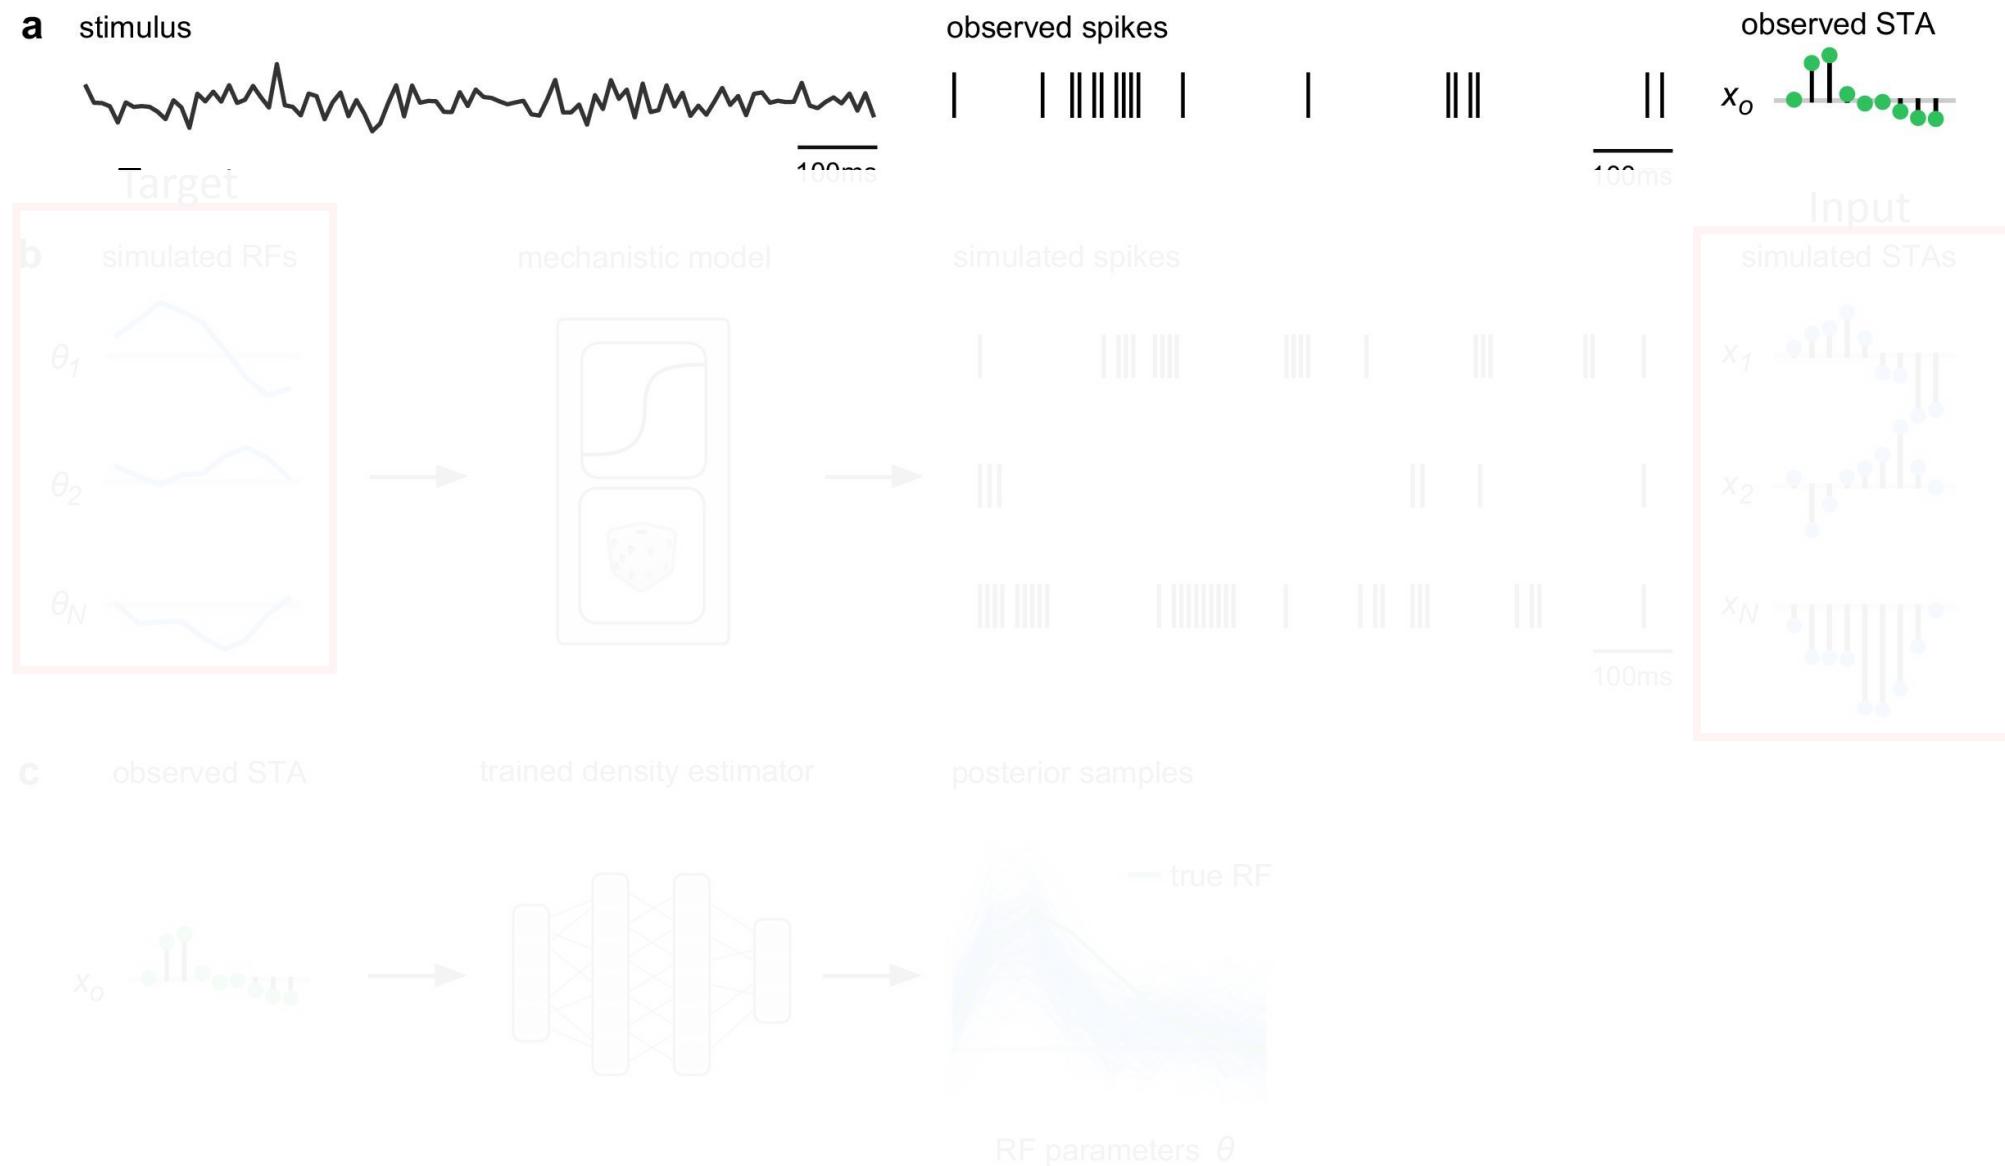
### Posterior Estimation



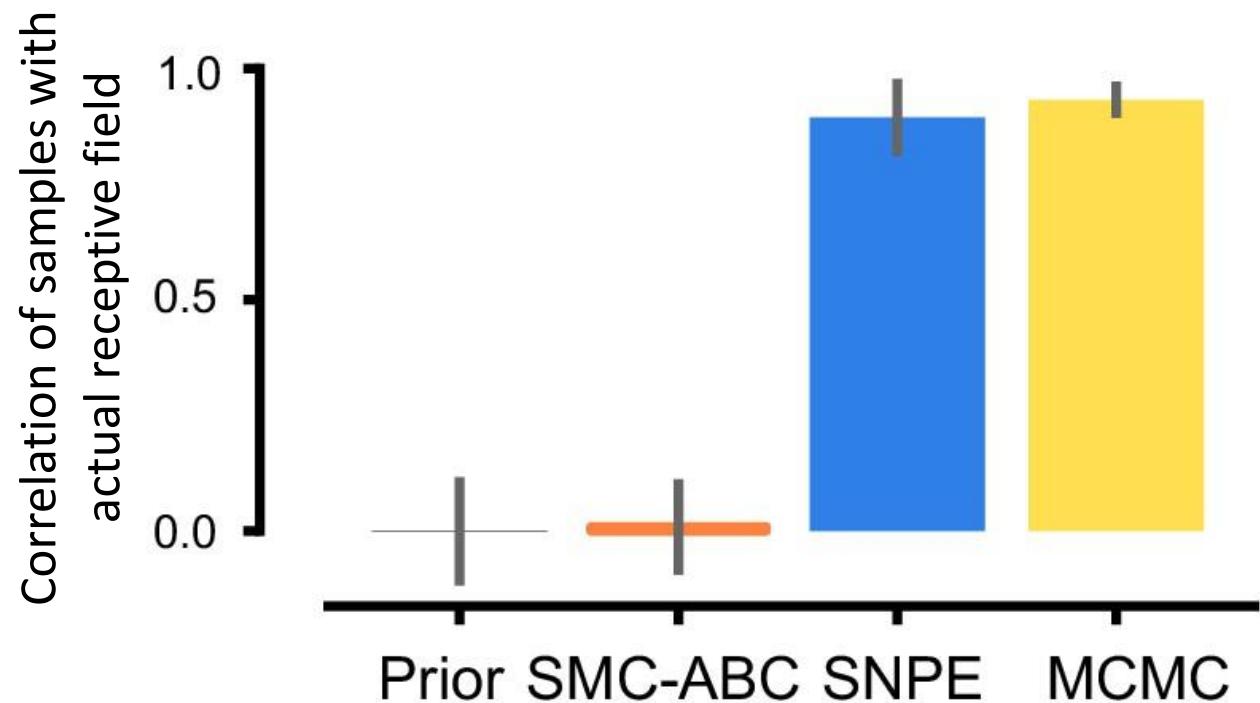
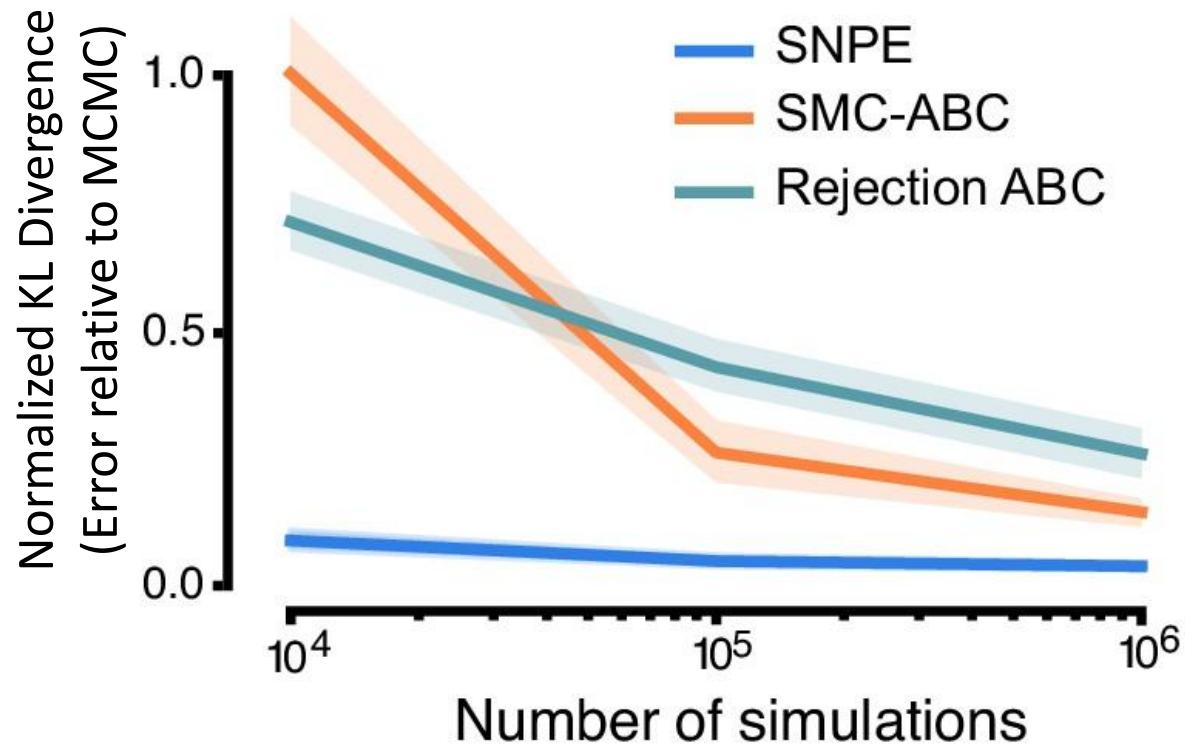
# Similar “Interesting” Example



# Results - Inferring receptive field (Linear-Nonlinear encoding model)



# Results - Comparison with established methodologies (Inferring receptive field)

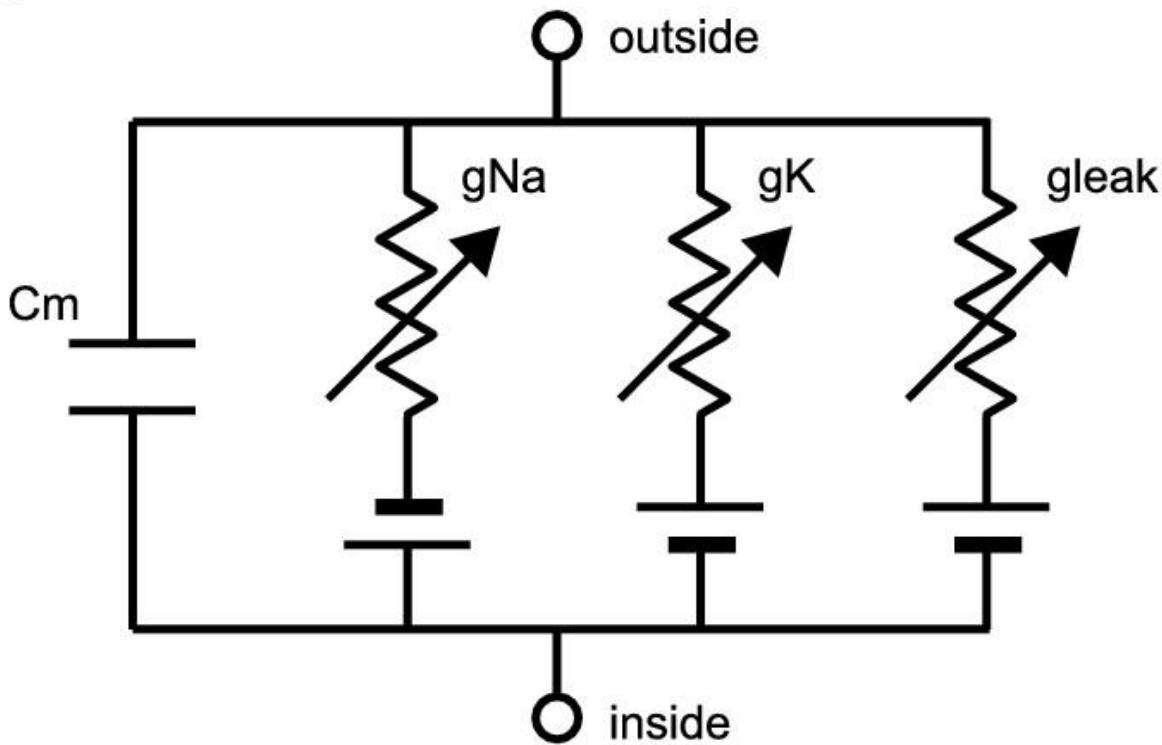


# Layout

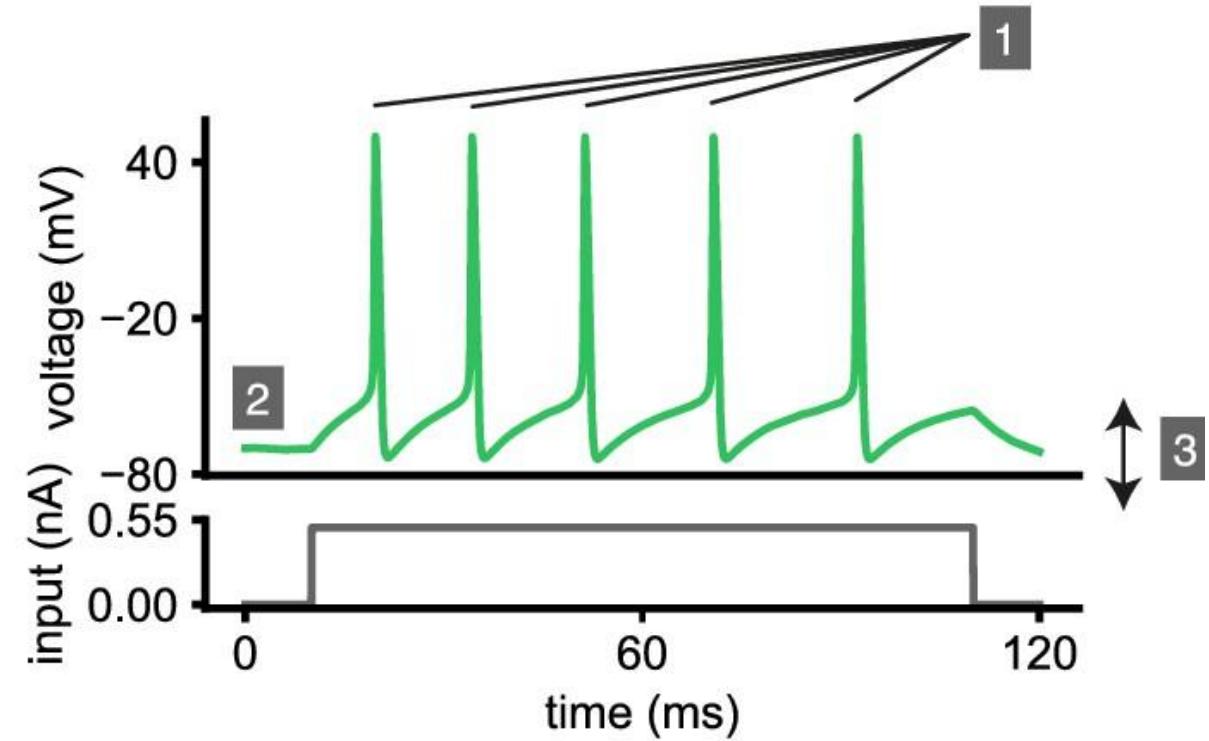
- Motivation
- Caveats (what this talk will tell you, and what it won't)
- Background
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - Related work (epsilon / likelihood-free methods)
    - ABC and SMB-ABC
- **Results**
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. ~~Inferring ion channel parameters~~
  5. **Inferring parameters for Hodgkin-Huxley Model**
  6. Inferring parameters for STG and analyzing posterior distribution

# Results - Inferring parameters for Hodgkin-Huxley Model

a



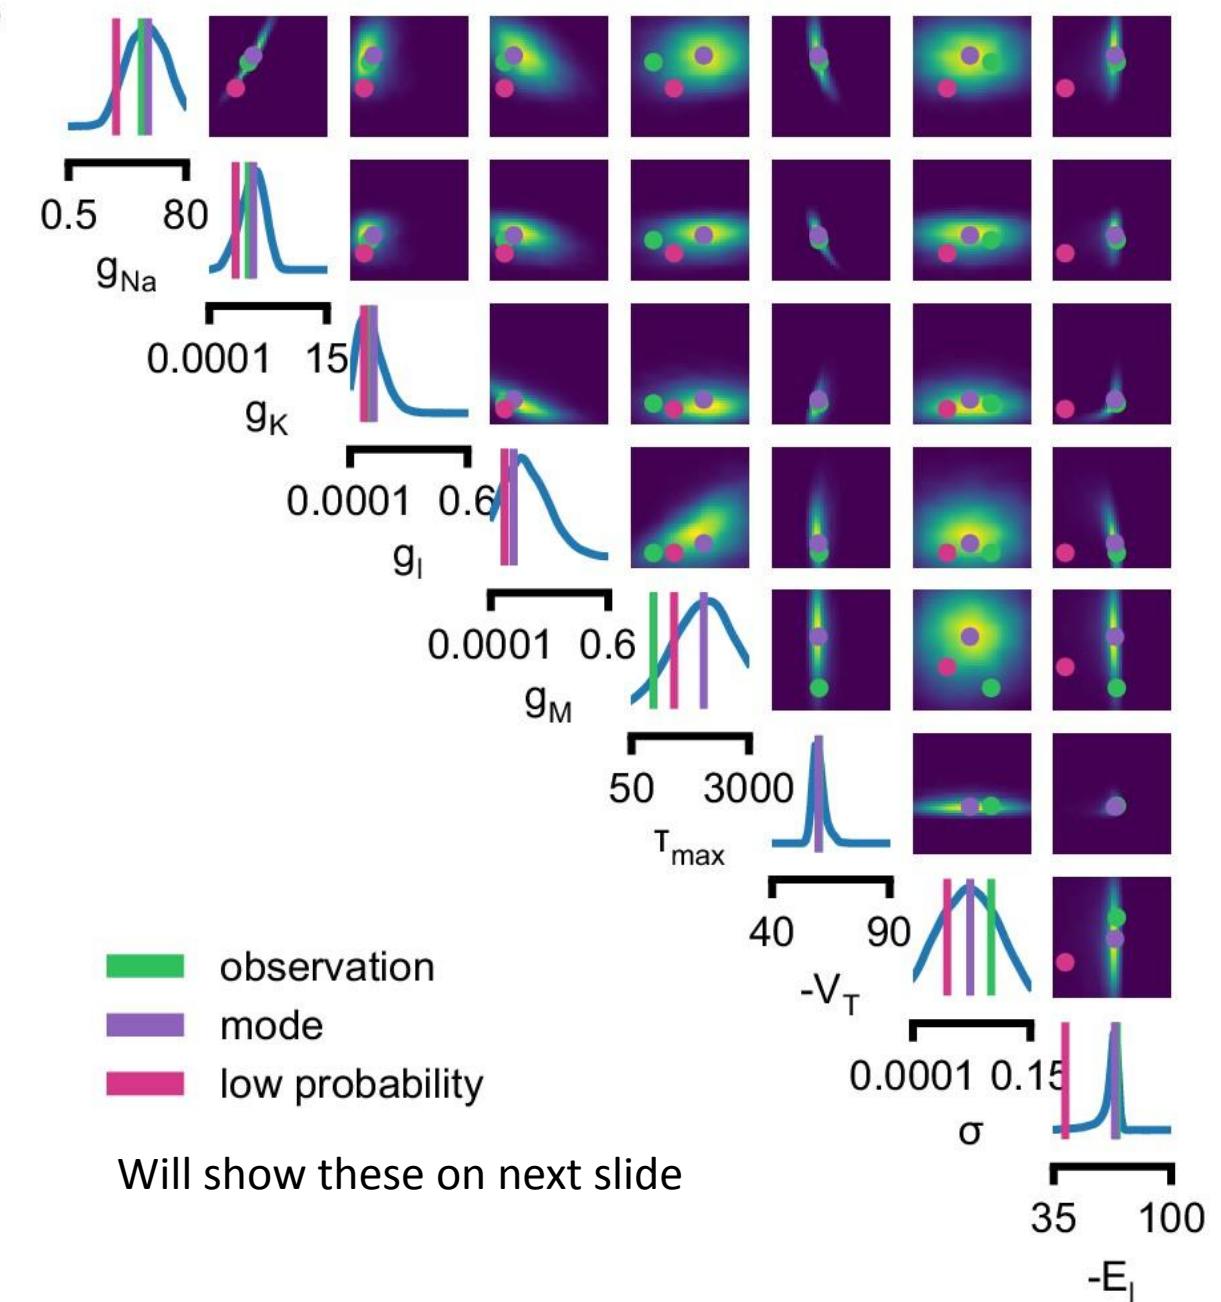
- (1) number of spikes,
- (2) mean resting potential
- (3) standard deviation of the pre-stimulus resting potential.



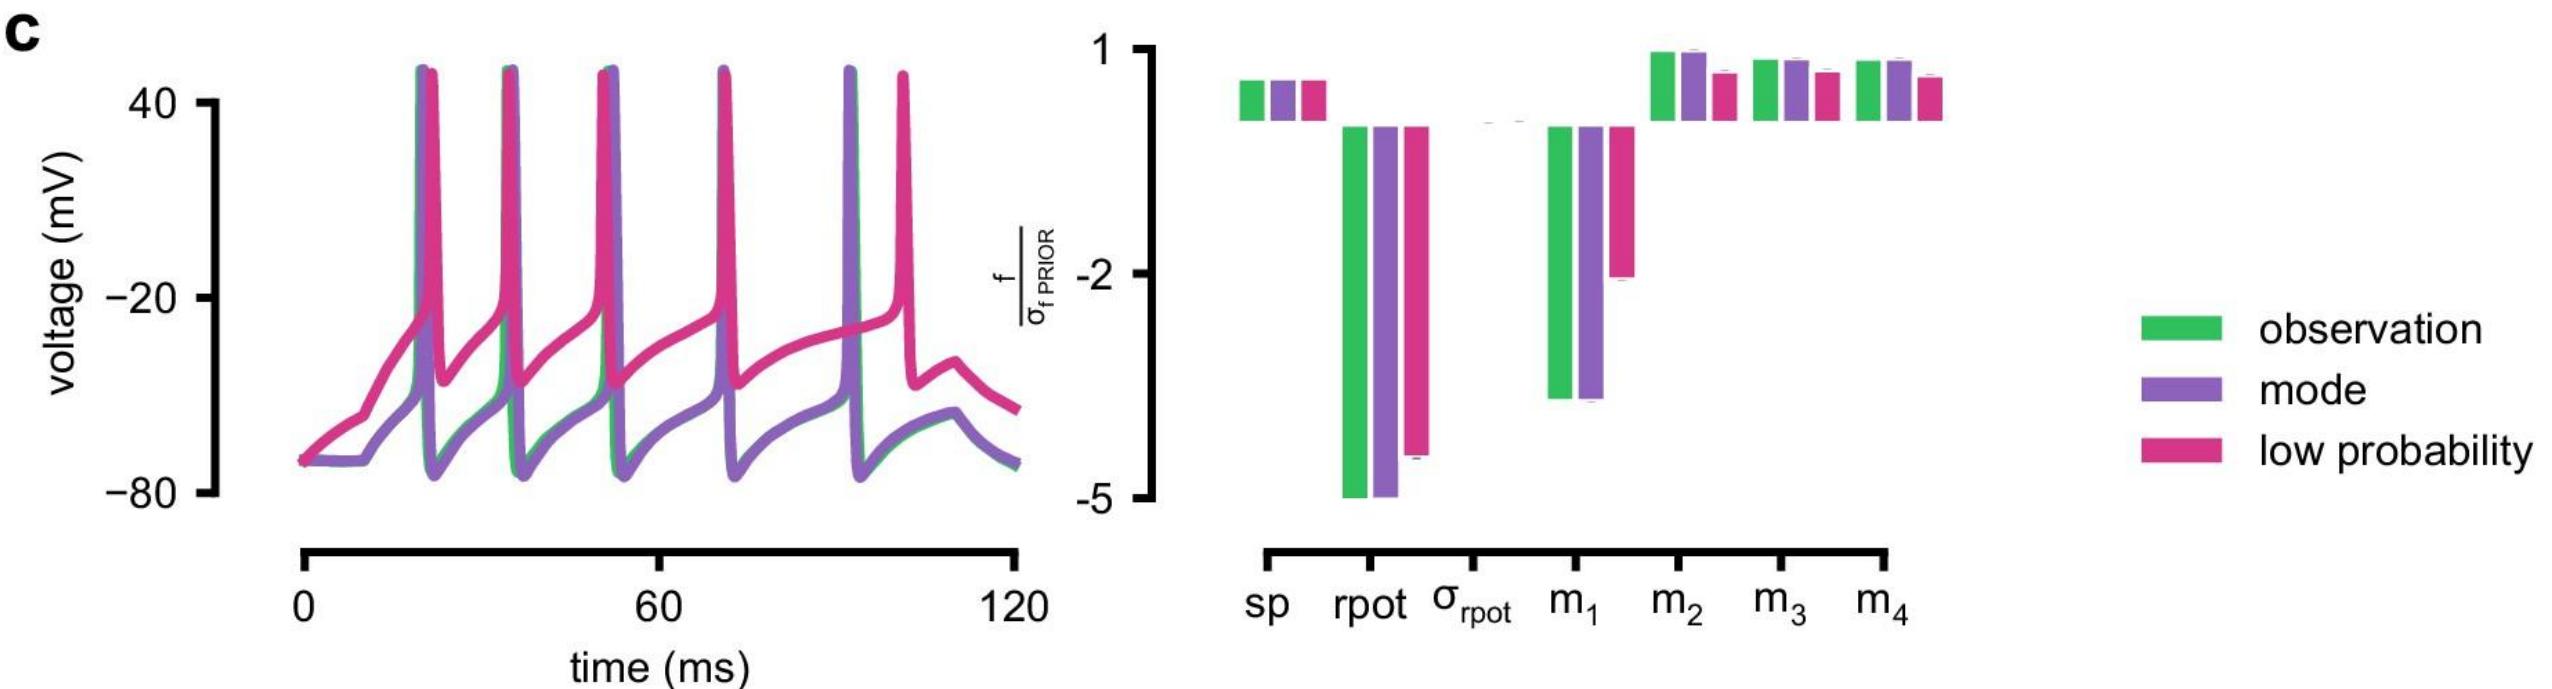
100,000 samples for 8 parameters  
 $5^8 \approx 400,000$

# Results - Inferring parameters for Hodgkin-Huxley Model

- Interesting to note parameters which need to be tightly vs loosely tuned
- **Importantly :** Parameters for observation are NOT mode of posterior



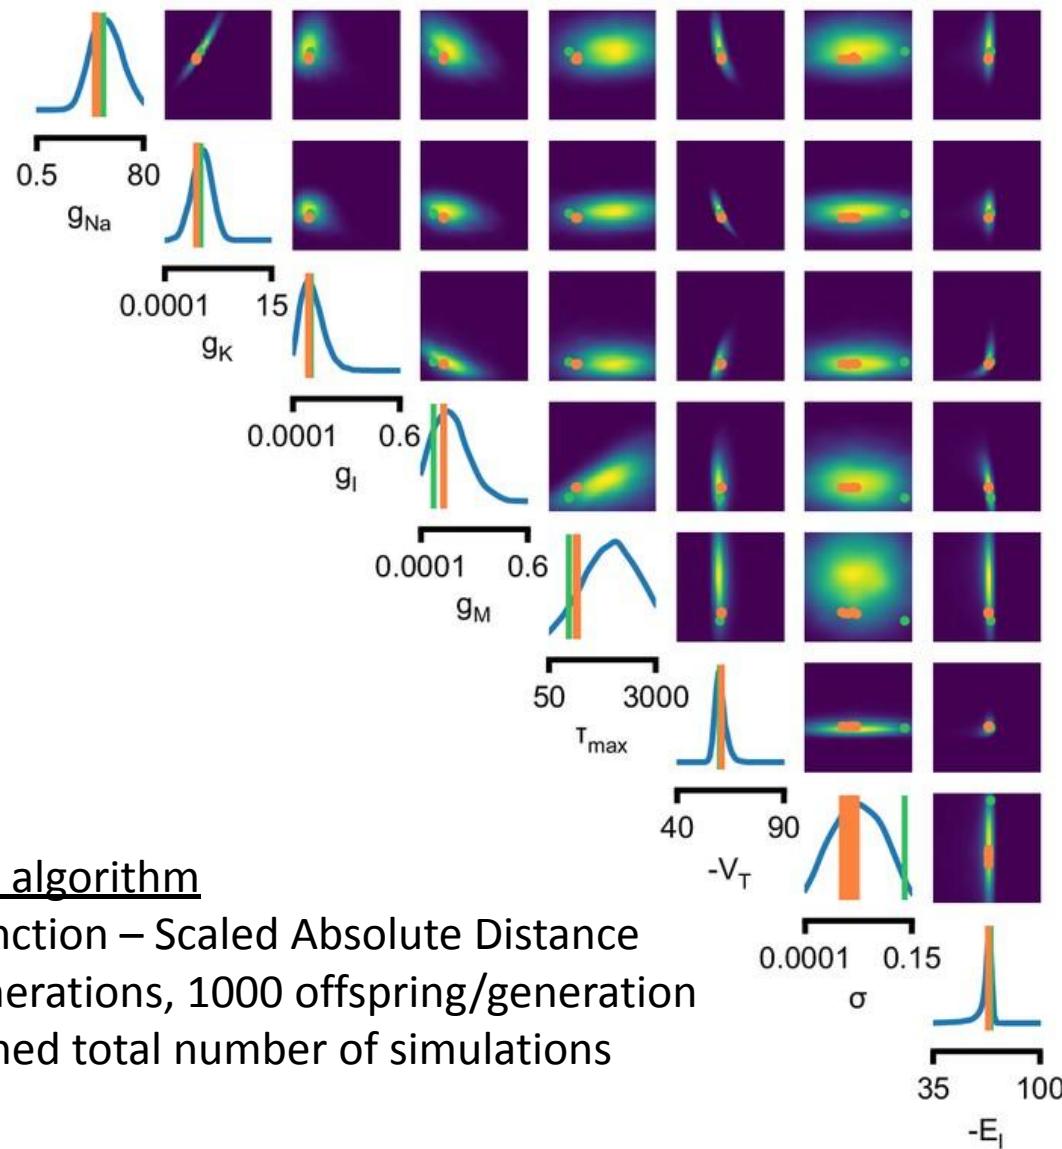
# Results - Inferring parameters for Hodgkin-Huxley Model



sp : number of spikes,  
rpot : mean resting potential,  
 $\sigma_{\text{rpot}}$  : standard deviation of the resting potential,  
first four voltage moments: mean ( $m_1$ ), standard deviation ( $m_2$ ), skewness ( $m_3$ ) and kurtosis ( $m_4$ )

# Results - Comparison with genetic algorithm

a



b

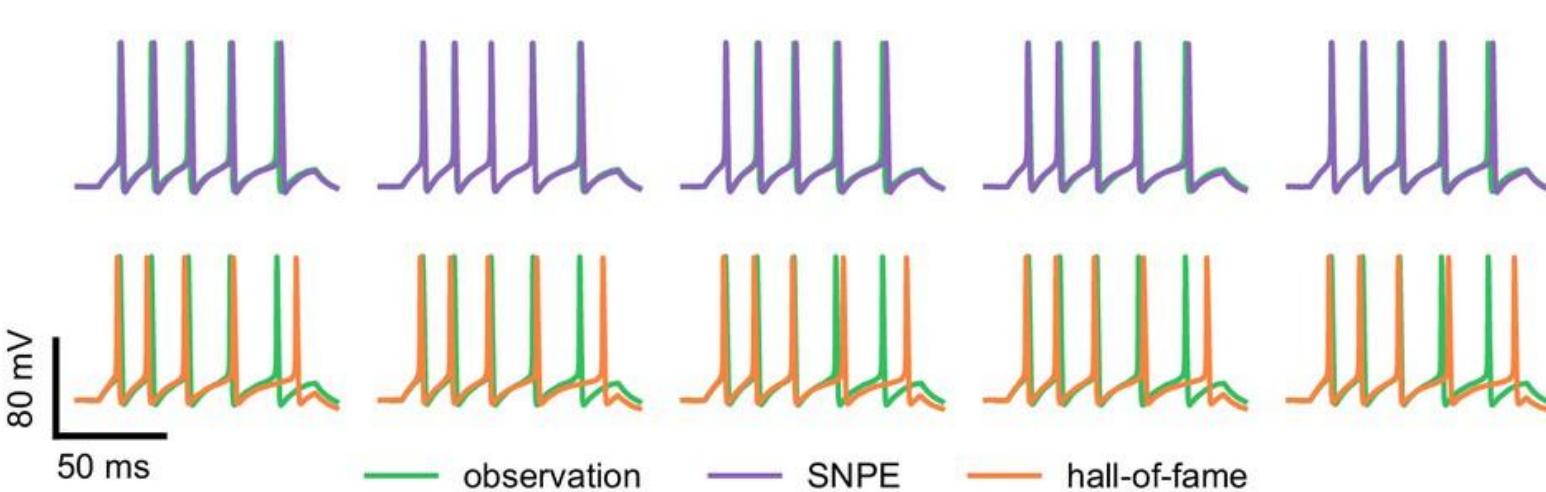


## Genetic algorithm

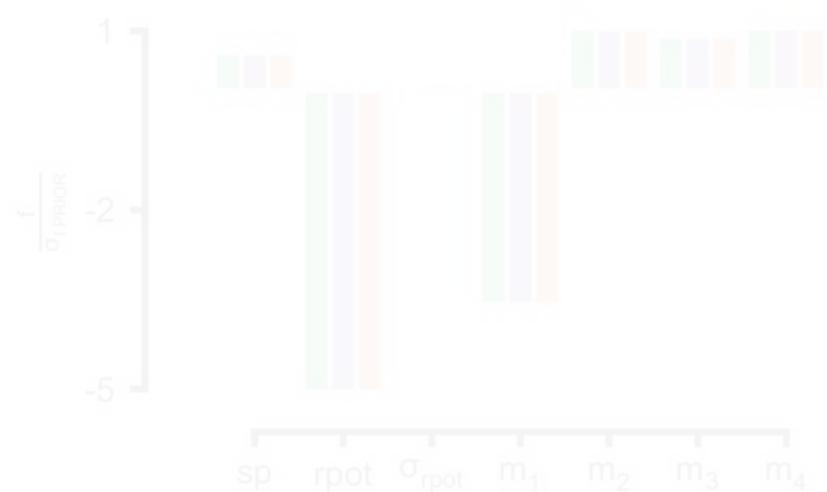
Cost function – Scaled Absolute Distance  
100 generations, 1000 offspring/generation  
 Matched total number of simulations

# Results - Comparison with genetic algorithm

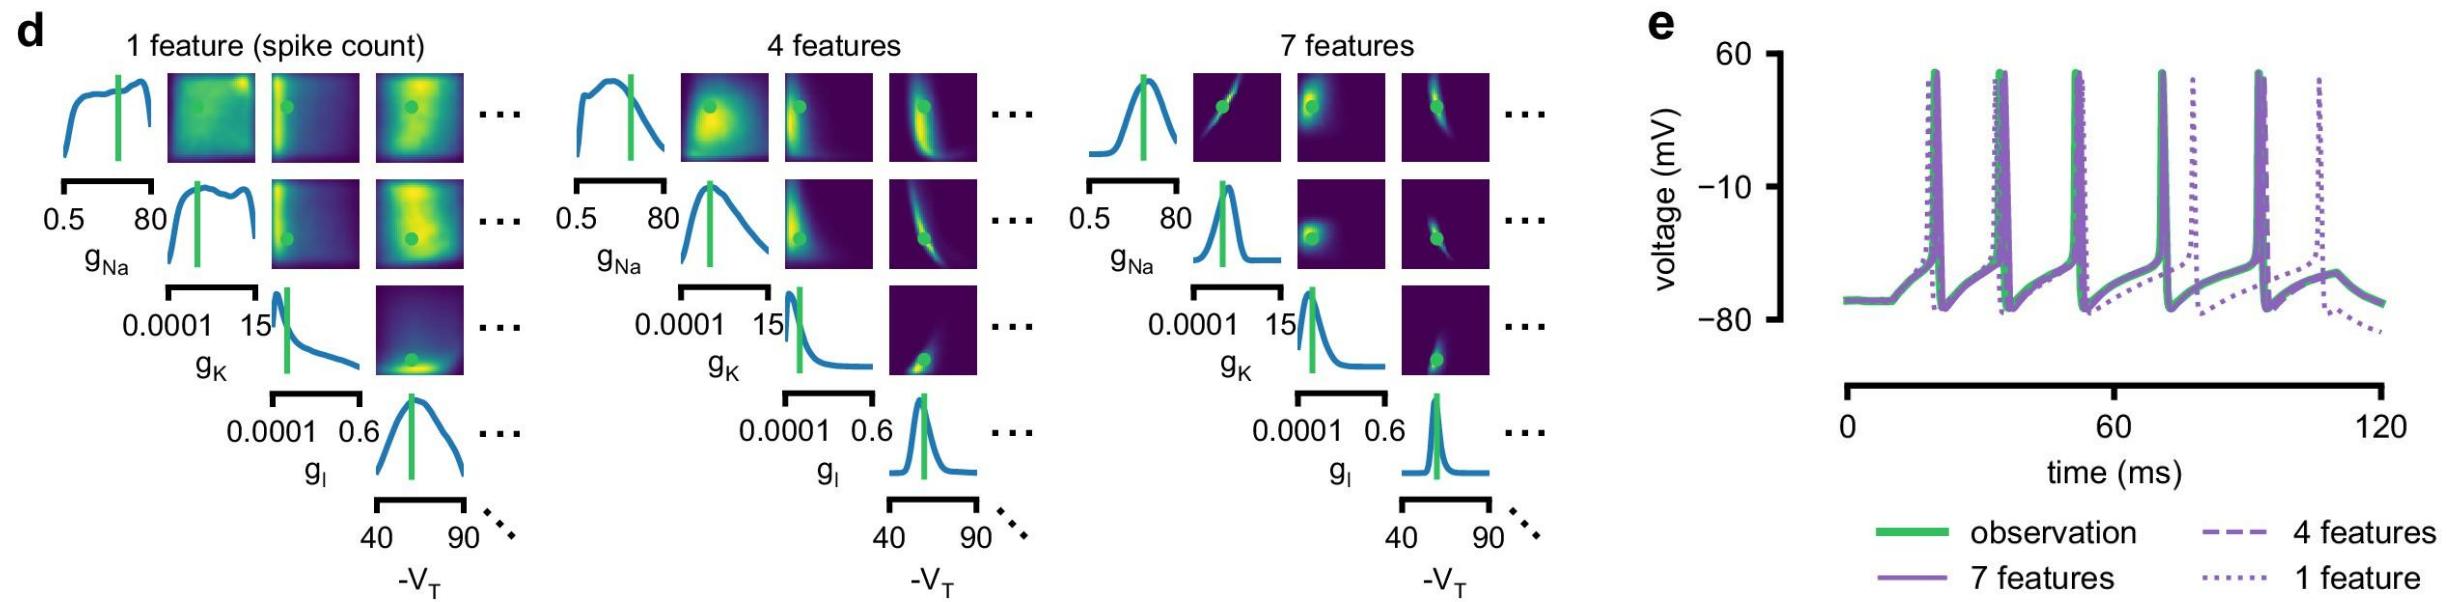
c



d



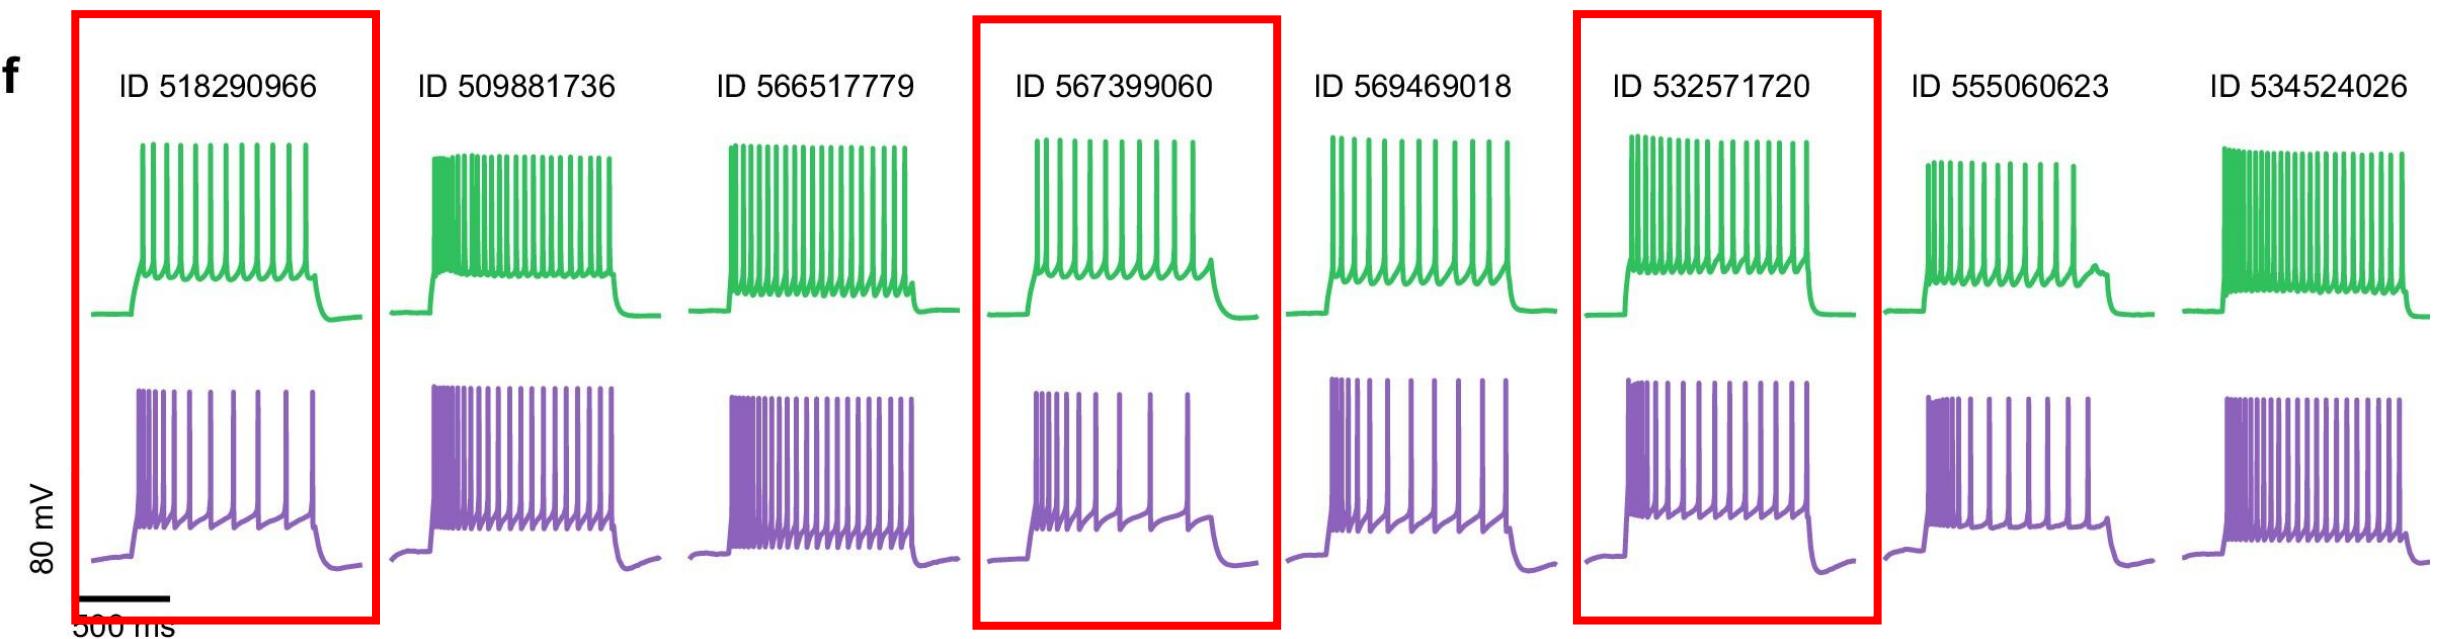
# Results - Inferring parameters for Hodgkin-Huxley Model



The more features you use, the tighter the distribution becomes

# Results - Inferring parameters for Hodgkin-Huxley Model

Observations from Allen Cell Types Database (green) and corresponding mode samples (purple)



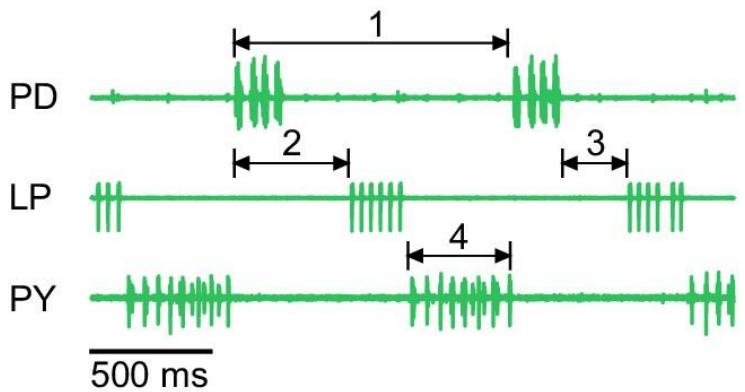
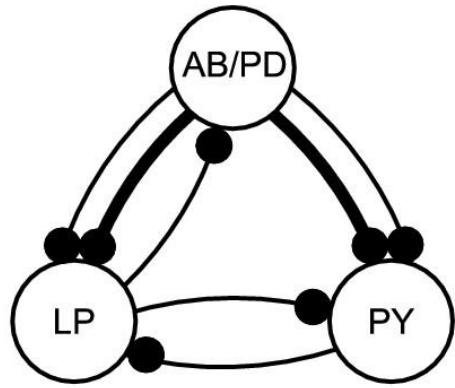
- Poor fits suggest amortized inference claim likely needs to be tempered
- The authors note that some parameters match better than others

# Layout

- Motivation
- Caveats (what this talk will tell you, and what it won't)
- Background
  - Conditional Density Estimation
  - Mixing Neural Networks with Probabilistic Models
  - Related work (epsilon / likelihood-free methods)
    - ABC and SMB-ABC
- **Results**
  1. Recap of Methodology
  2. Comparison with established methodologies
  3. Inferring visual receptive field
  4. ~~Inferring ion channel parameters~~
  5. Inferring parameters for Hodgkin-Huxley Model
  6. **Inferring parameters for STG and analyzing posterior distribution**

# Results - Inferring parameters for STG

a

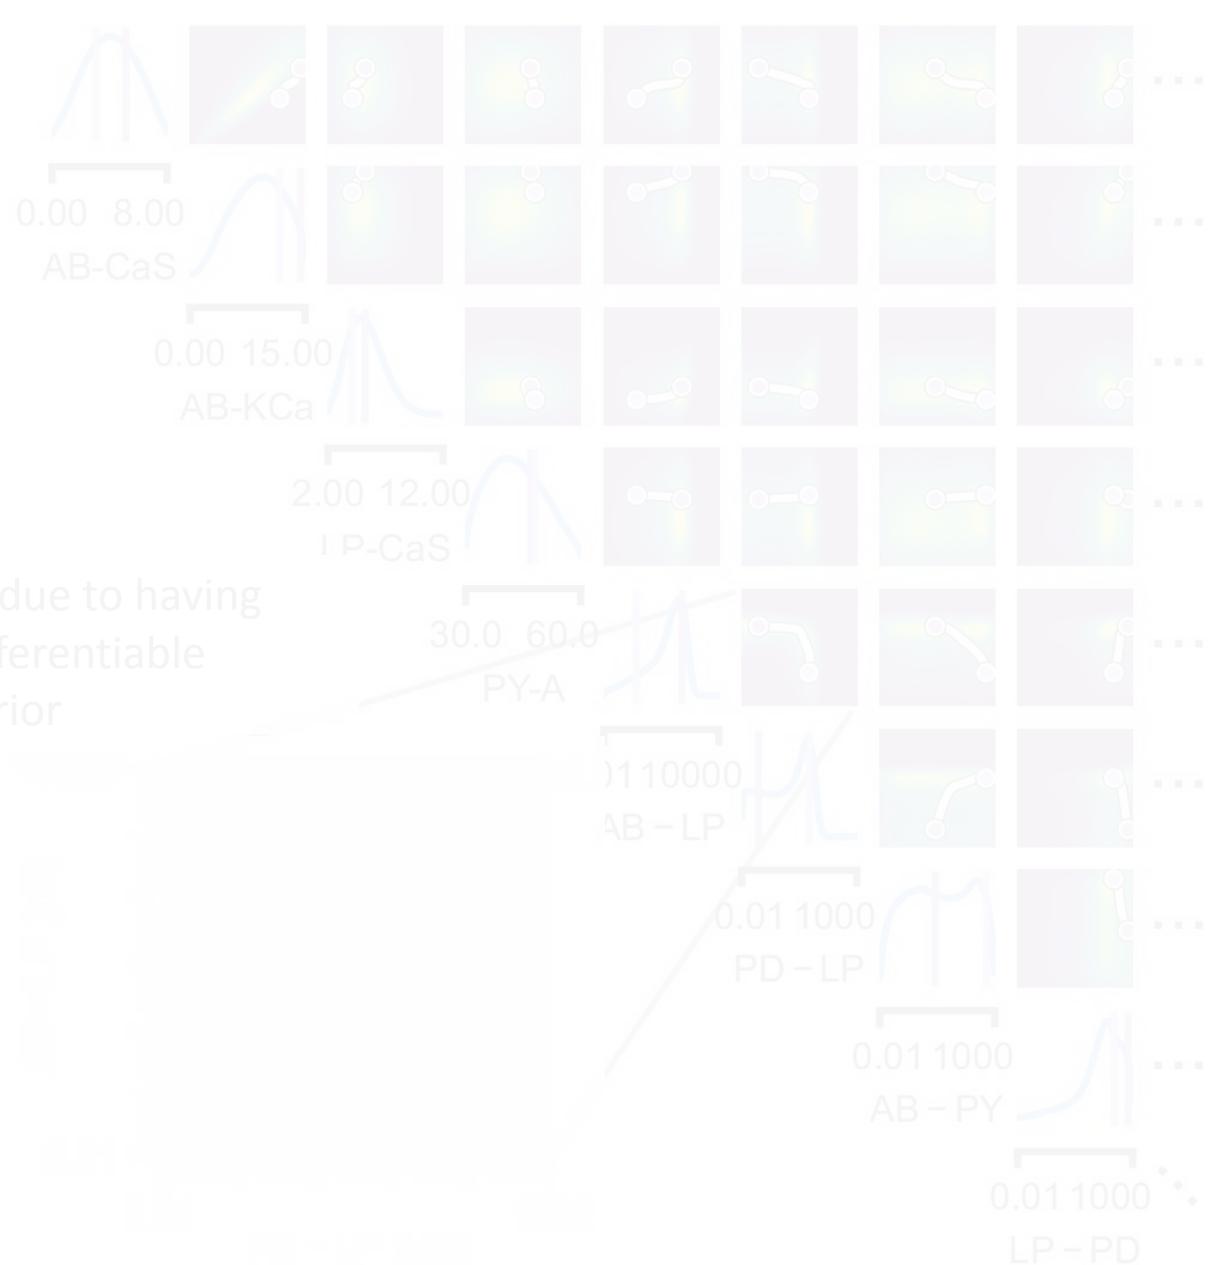


1. cycle period
2. phase delays
3. phase gaps
4. burst durations

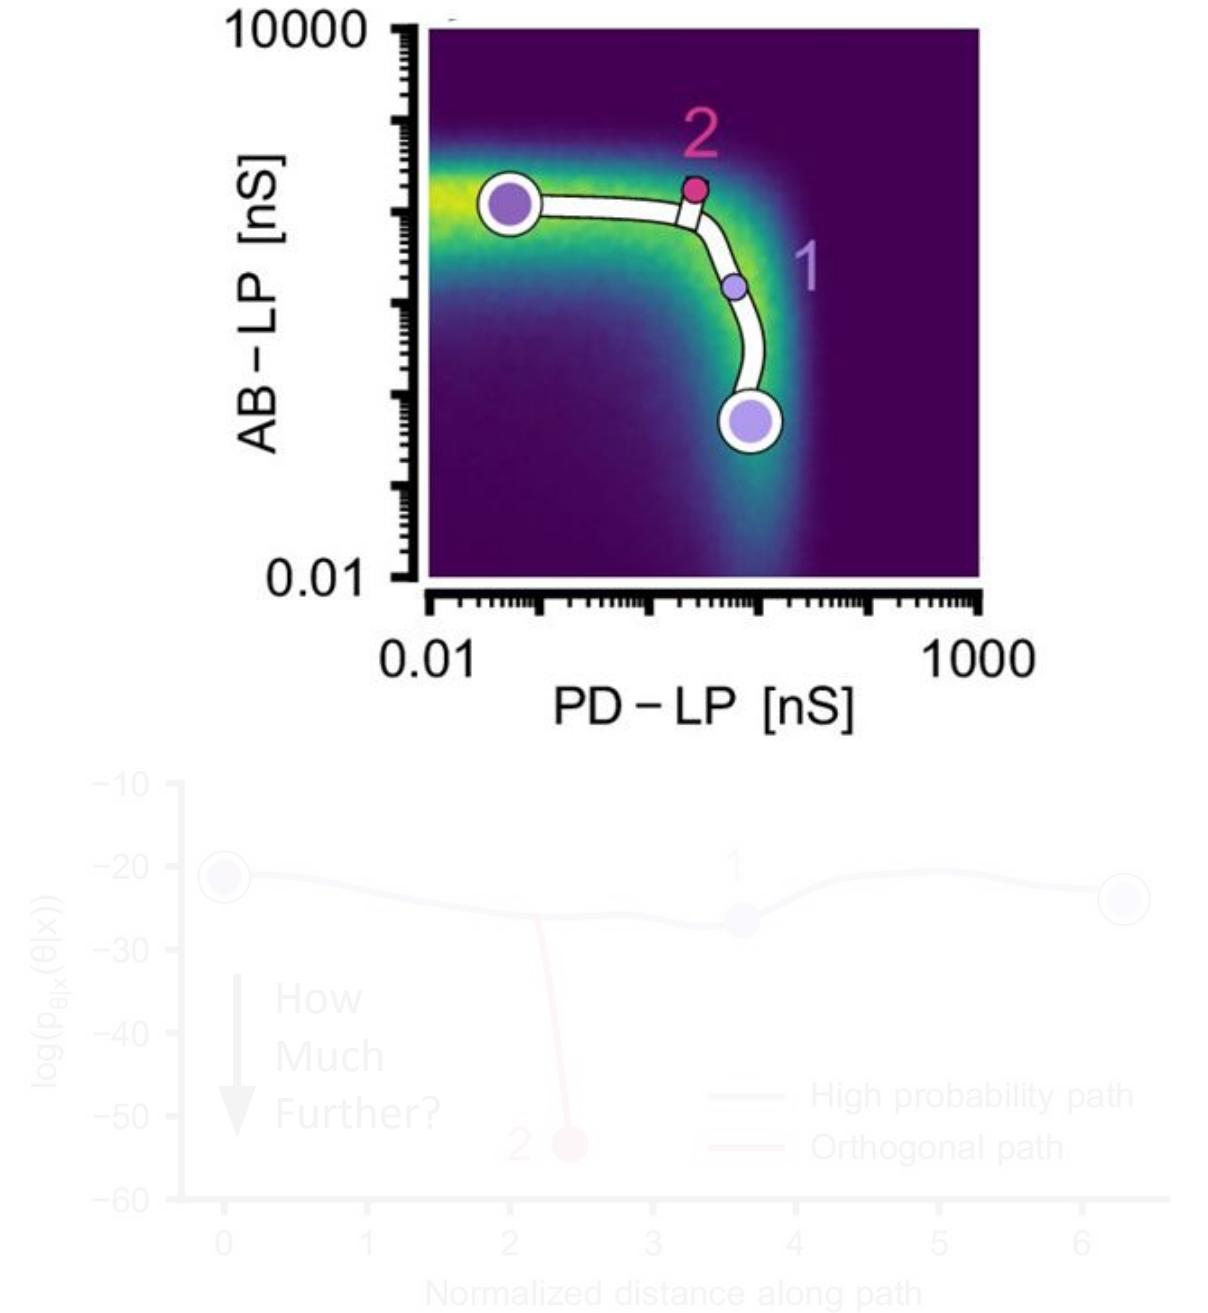
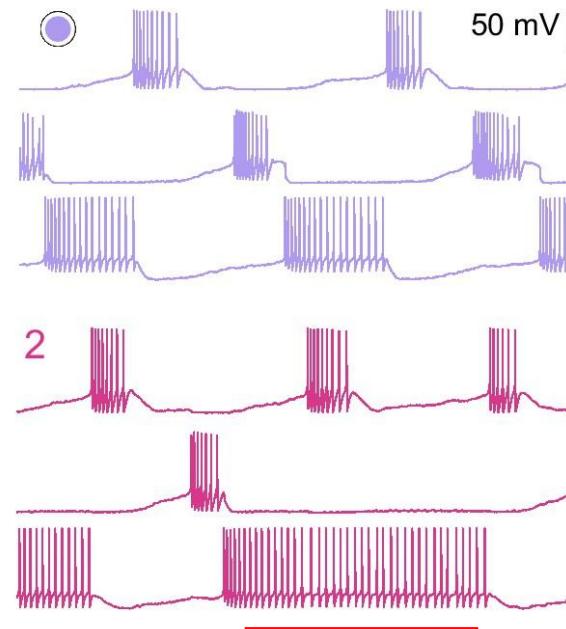
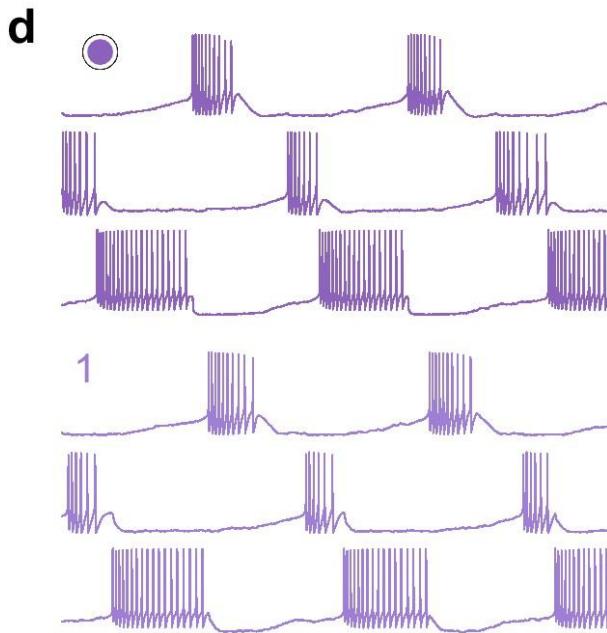
24 membrane parameters and 7 synaptic parameters

31 parameters in total

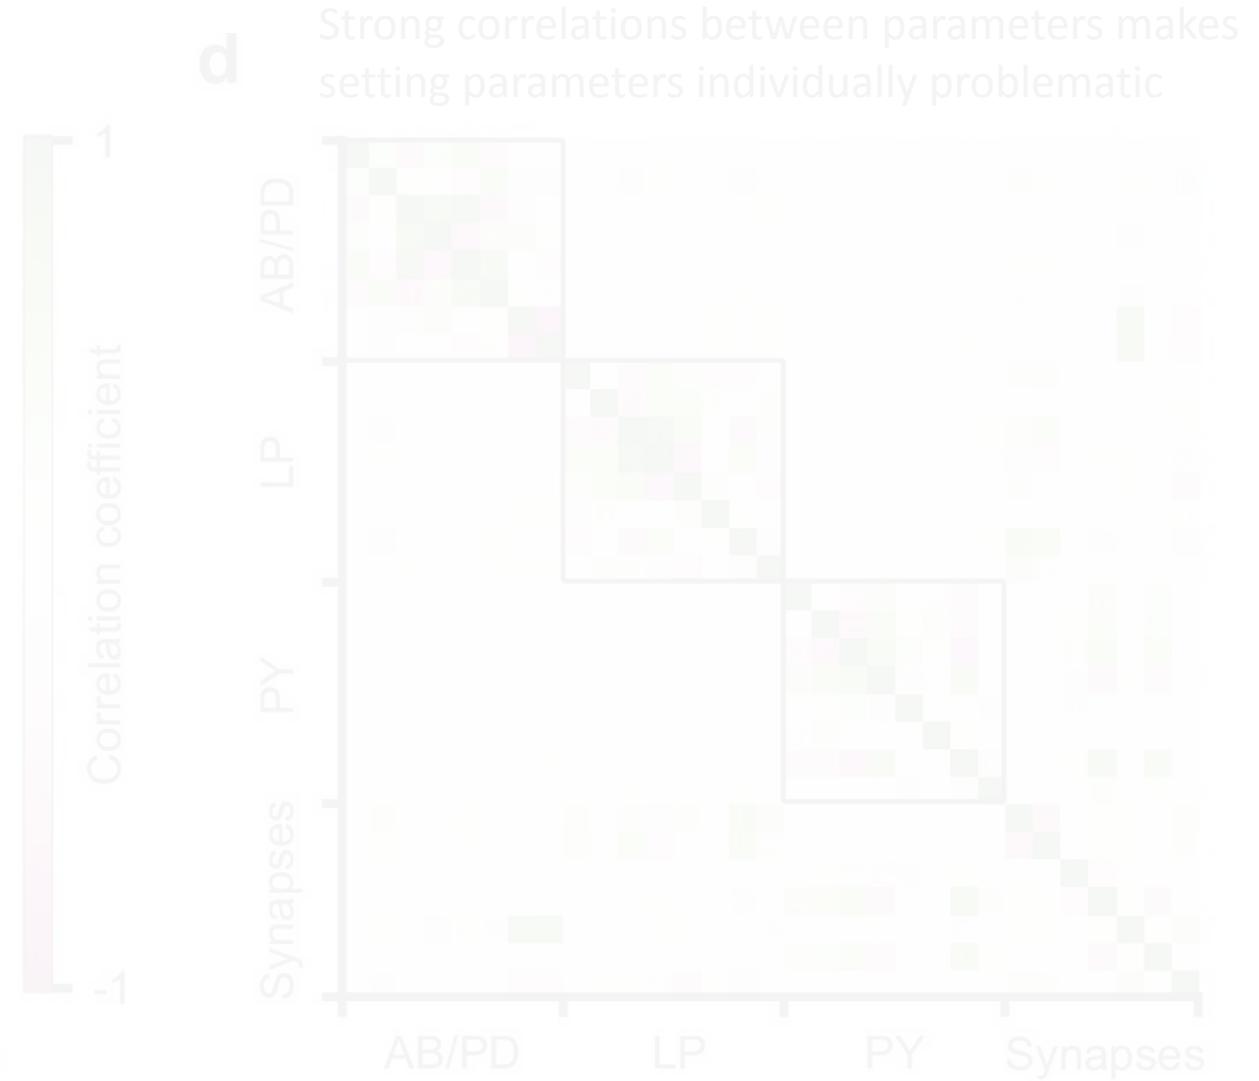
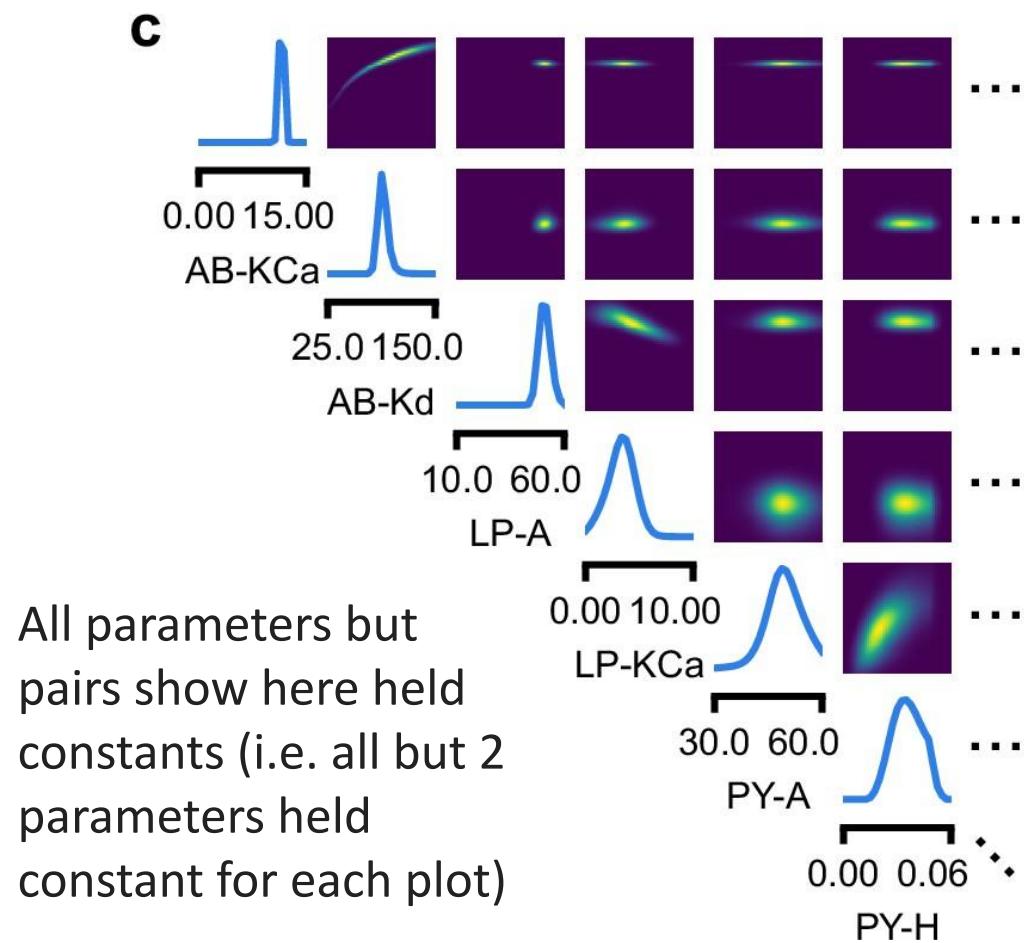
18 million samples used here vs.  $2^{31}$  (~2 billion) for grid with 2 values per parameter



# Results - Inferring parameters for STG

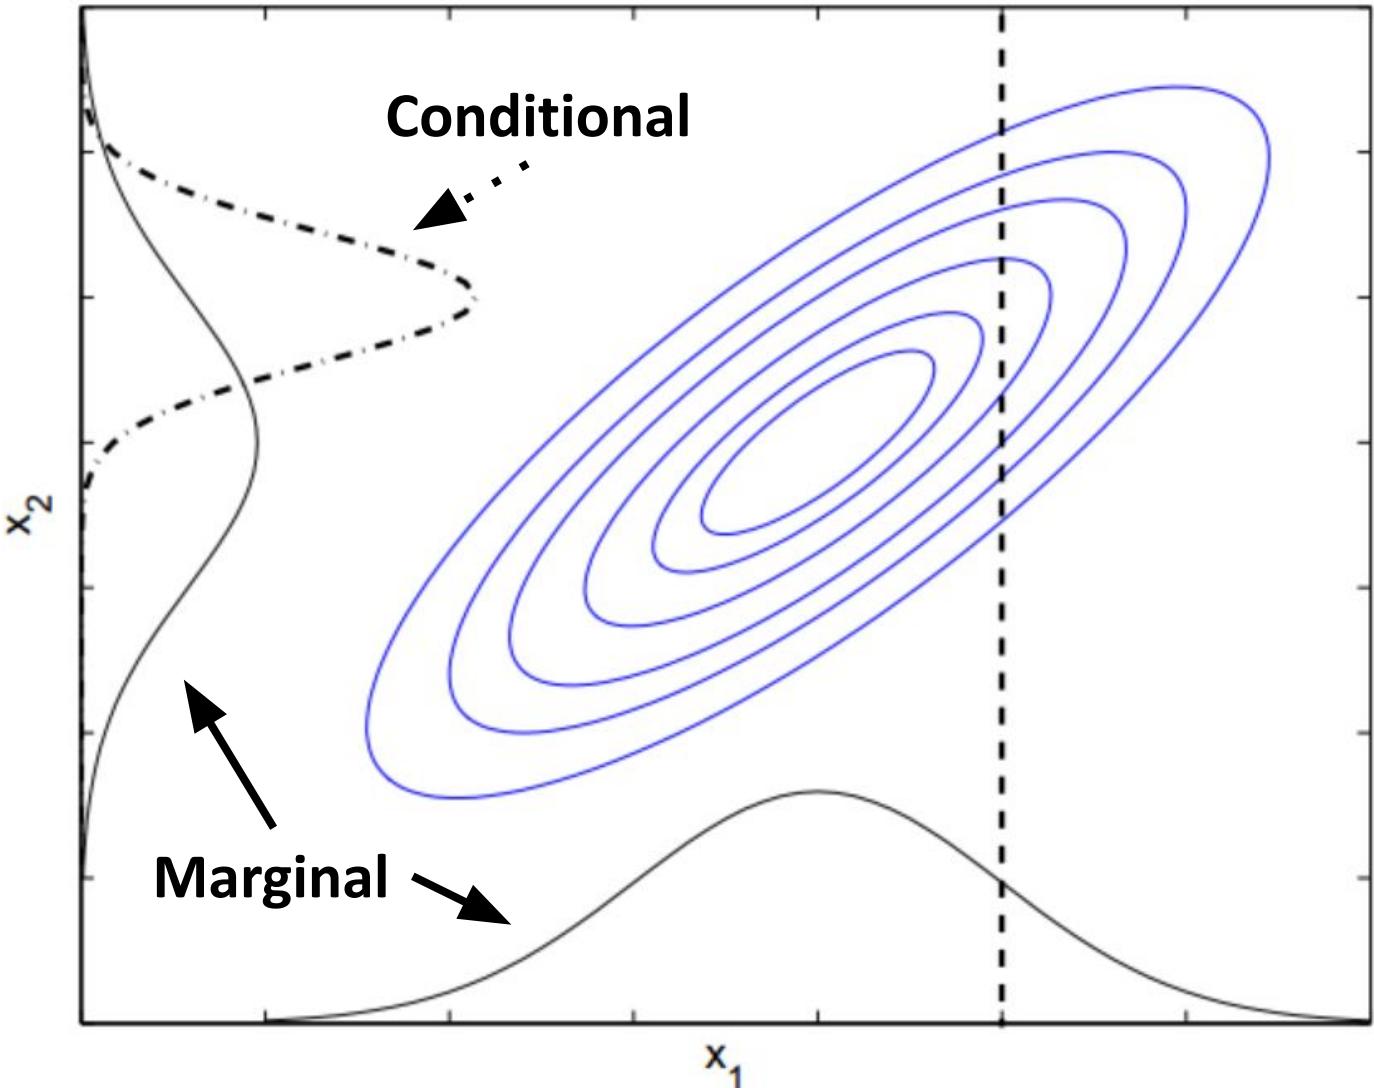


# Results – Analysis of STG Parameter Posterior



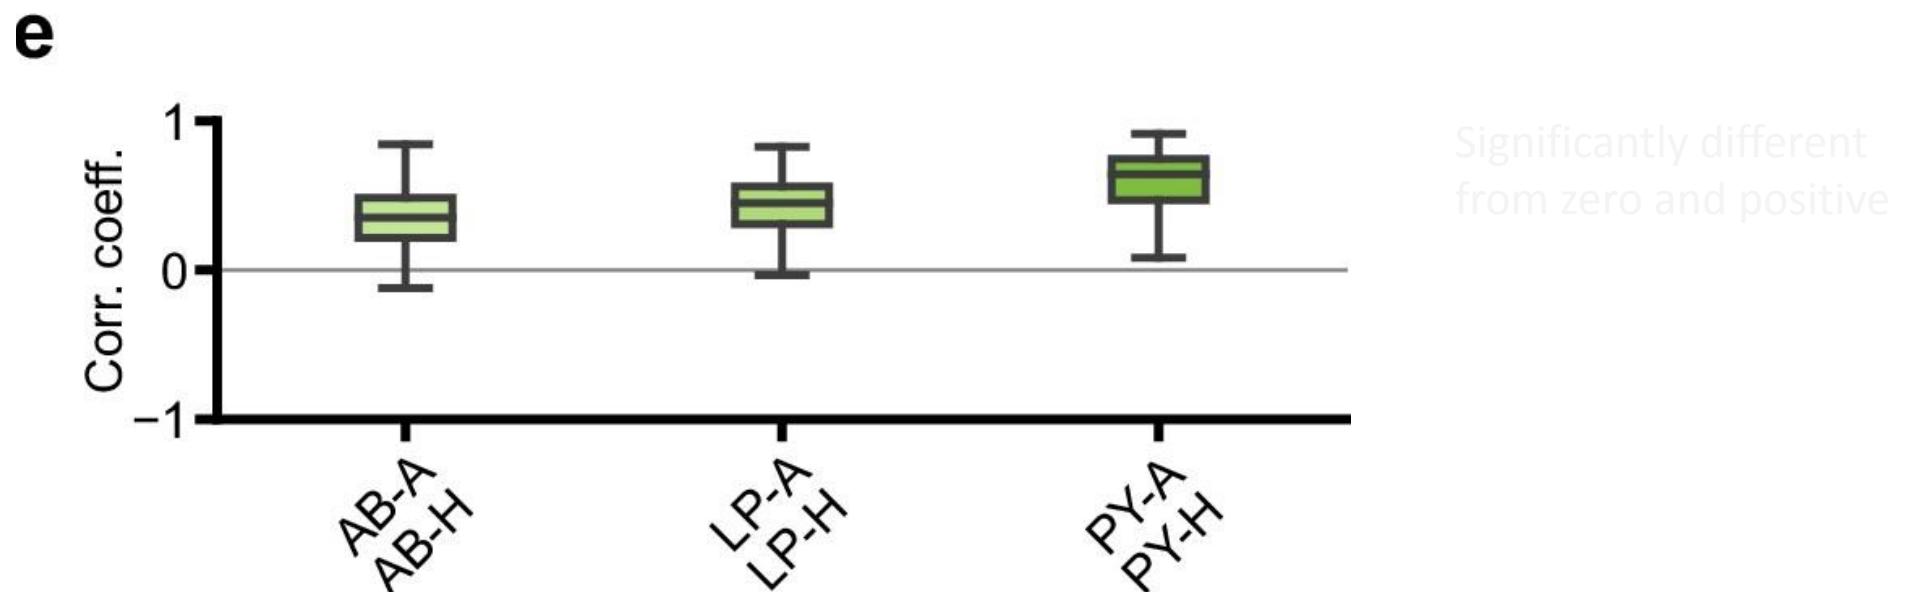
# Conditional vs Marginal Densities

- Marginal
  - Summing over (or ignoring) other variables
- Conditional
  - Holding other variables at constant value



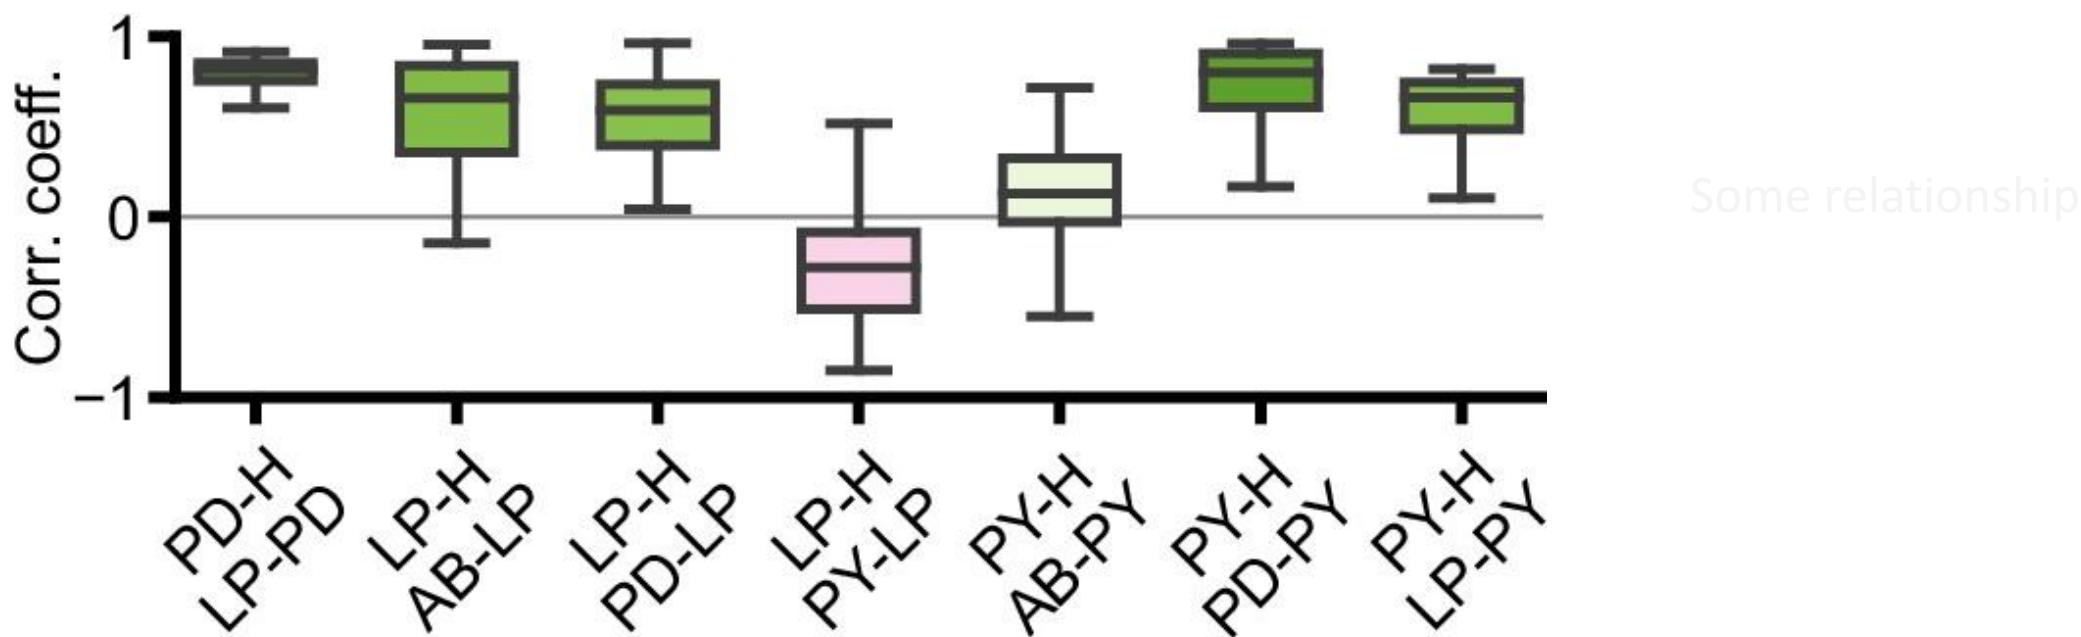
# Results – Analysis of STG Parameter Posterior

- For the PD and the LP neuron
- overexpression of the fast transient potassium current (IA) leads to a compensating increase of the hyperpolarization current (IH)
- (MacLean et al., 2003; MacLean et al., 2005).



# Results – Analysis of STG Parameter Posterior

- Using the dynamic clamp,
- Diverse combinations of the synaptic input strength and the maximal conductance of IH lead to similar activity in the LP and the PD neuron
- (Grashow et al., 2010; Marder, 2011).



# Conclusion

- Novel methodology to estimate posterior distribution of parameters for given data
  - Appears to be more efficient than current standard methods
  - Having full posterior allows exploring relationships between parameters
- 
- Criticism:
    - Did not compare with more advanced ABC samplers (ABC-MCMC) which might be more efficient
  - Advantage:
    - Amortization of inference (do not need to rerun inference for new sample once neural network has been trained)
    - Sequential inference (used, but not explored here) allows faster training
    - No need to define distance function or distance threshold
    - Uses all samples, unlike grid sampling

# Simulation based inference in Python

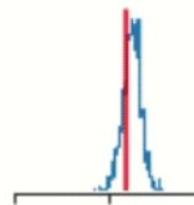
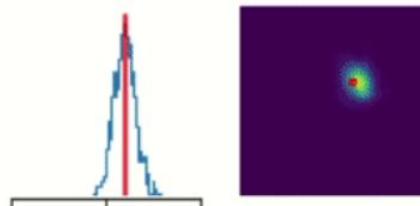
sbi: simulation-based inference

<https://www.mackelab.org/sbi/>

```
: prior = BoxUniform(low=zeros(2), high=2*ones(2)) # Box prior [0,2]x[0,2]
def simulator(theta): return theta + 0.1*randn_like(theta) # Gaussian in 2D
posterior = infer(simulator, prior, method='SNPE', num_simulations=500)
```

```
Running 500 simulations.: 100%|██████████| 500/500 [00:00<00:00, 57141.55it/s]
Neural network successfully converged after 109 epochs.
```

```
: samples = posterior.sample((1000,), x=observed)
pairplot(samples, points=ground_truth, **plot_style);
```



# Simulation based inference in Python

ELFI - Engine for Likelihood-Free Inference

<https://elfi.readthedocs.io/en/latest/>

## Currently implemented LFI methods:

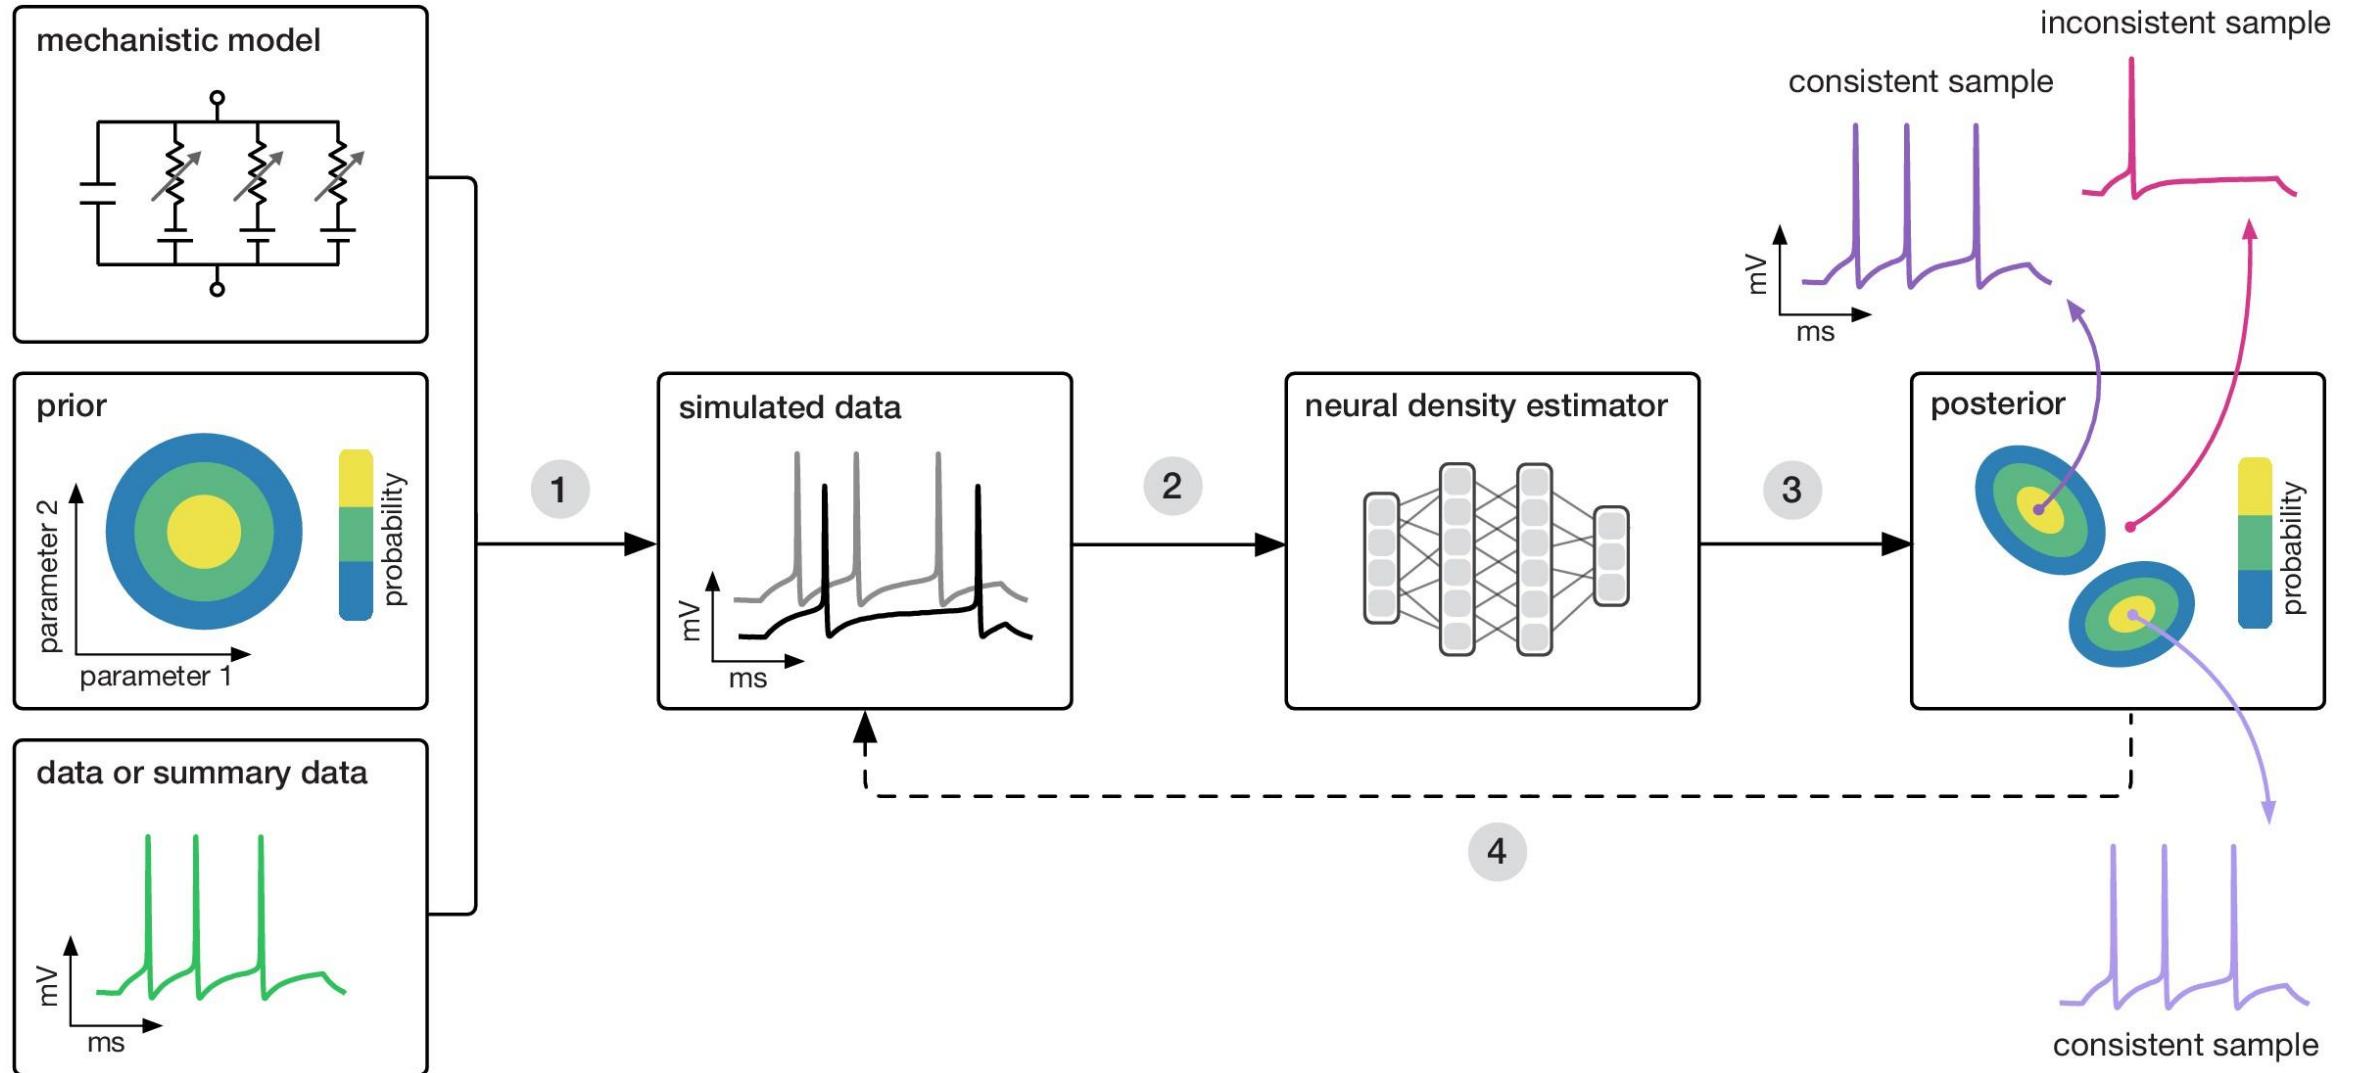
- ABC rejection sampler
- Sequential Monte Carlo ABC sampler
- ABC-SMC sampler with [adaptive distance](#)
- ABC-SMC sampler with [adaptive threshold selection](#)
- Bayesian Optimization for Likelihood-Free Inference (**BOLFI**) framework
- Robust Optimization Monte Carlo (**ROMC**) framework

ELFI also has the following non LFI methods:

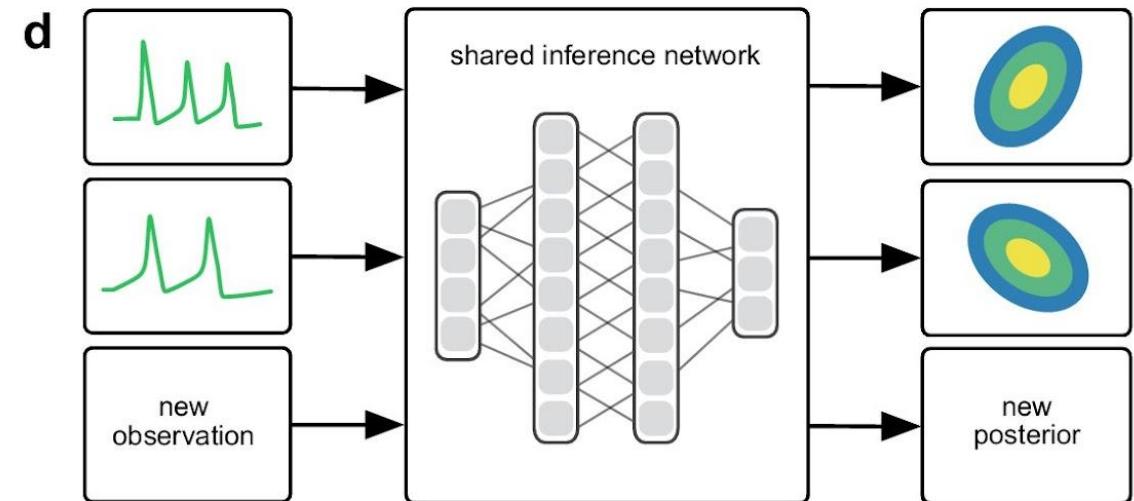
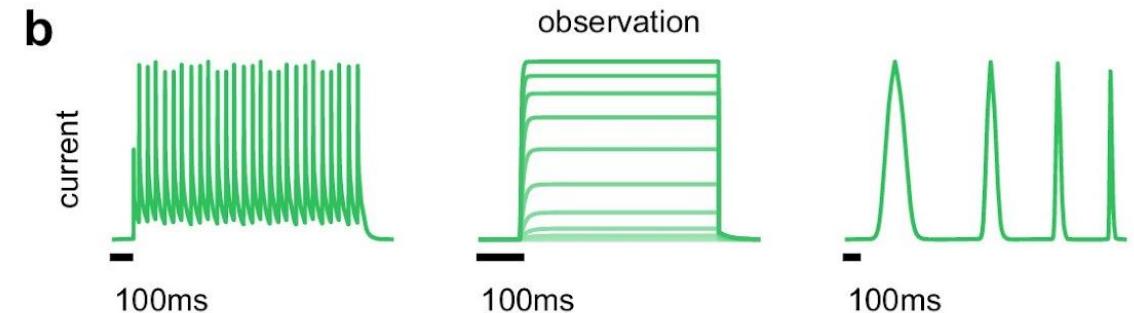
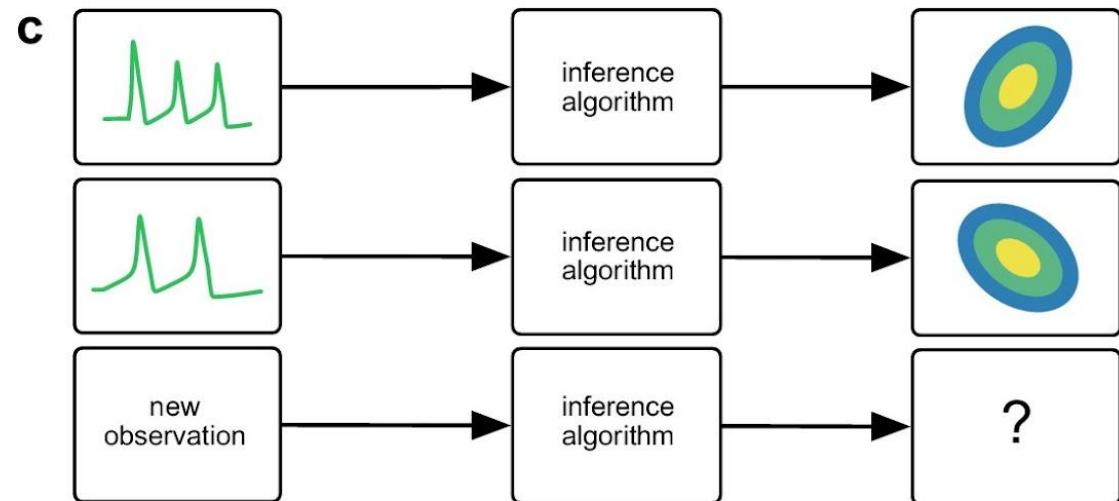
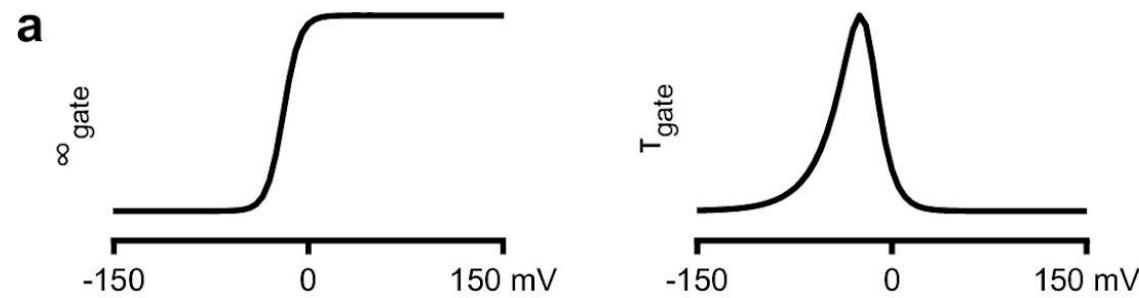
- Bayesian Optimization
- [No-U-Turn-Sampler](#), a Hamiltonian Monte Carlo MCMC sampler

Additionally, ELFI integrates tools for visualization, model comparison, diagnostics and post-processing.

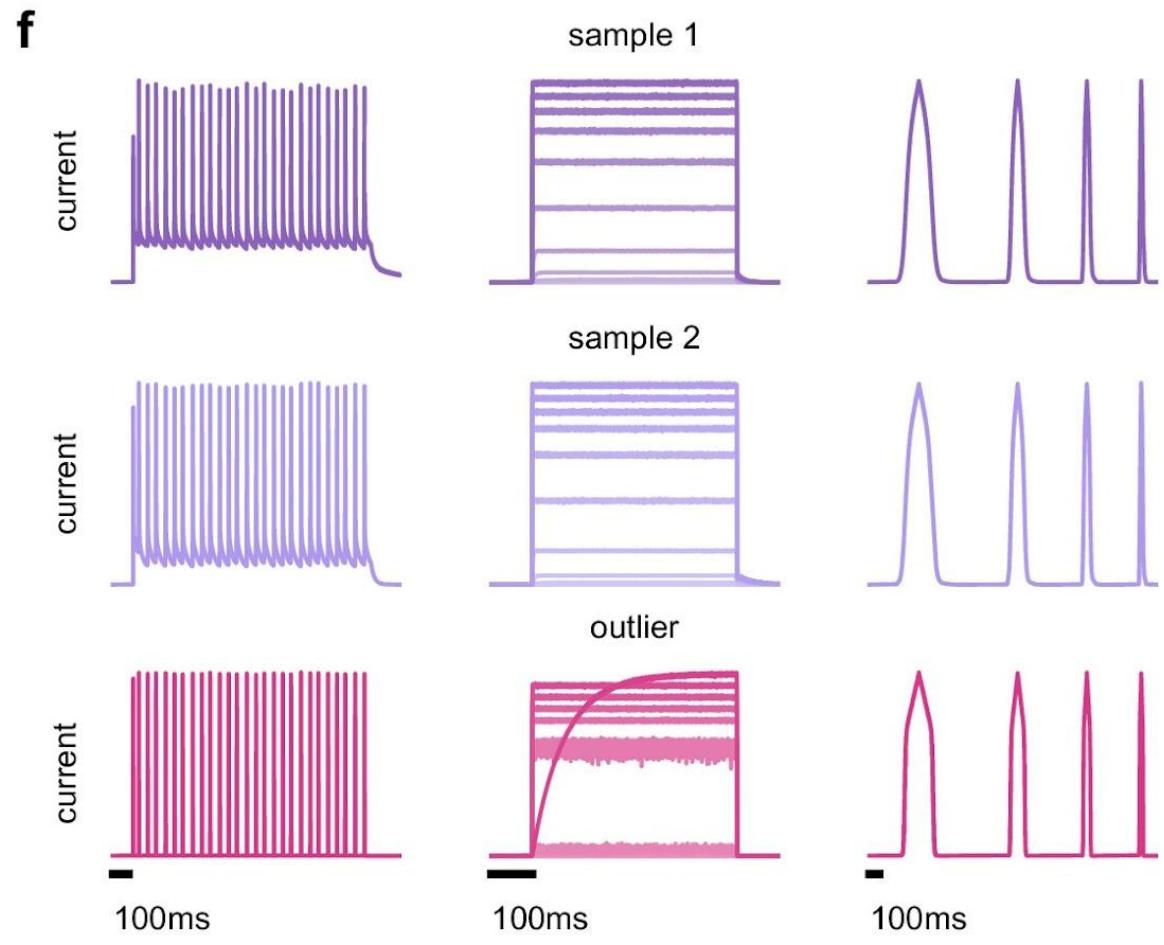
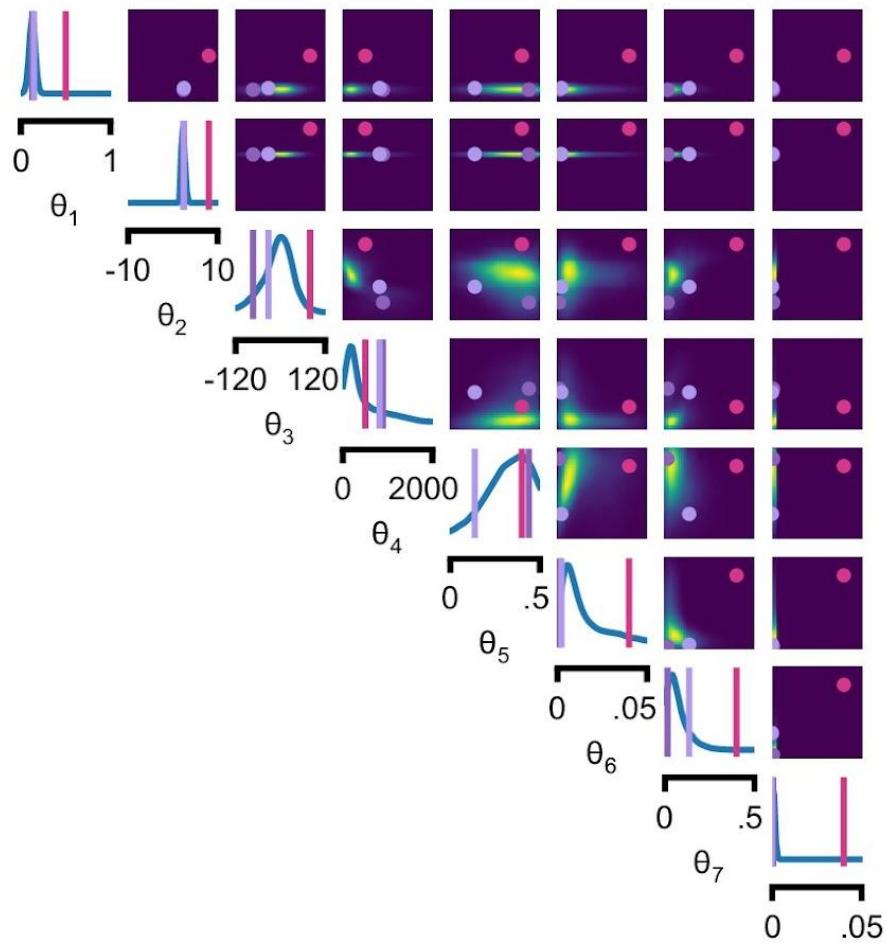
# Questions?



# Results - Inferring ion channel parameters



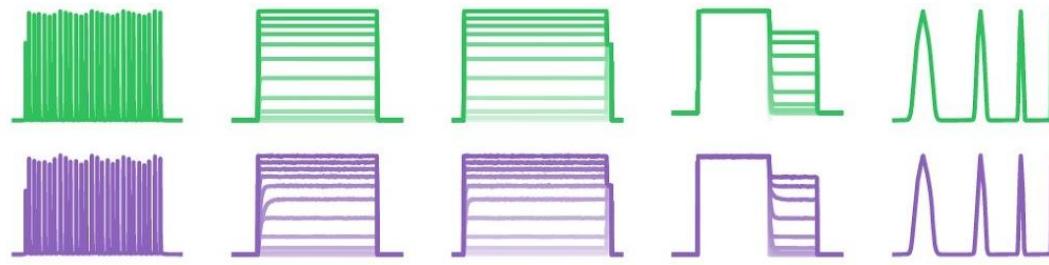
# Results - Inferring ion channel parameters



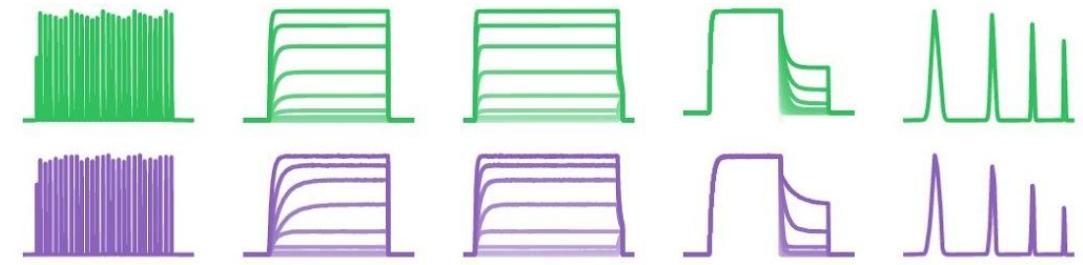
# Results - Inferring ion channel parameters

**g**

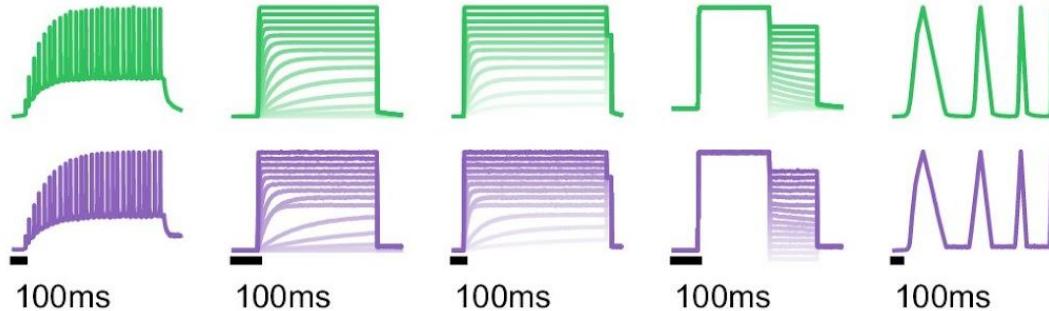
Channel 144482\_kv · CC=0.999



Channel 181967\_kdrca1 · CC=0.994



Channel 185858\_IKM · CC=0.997



Channel 116094\_kM · CC=0.974

